


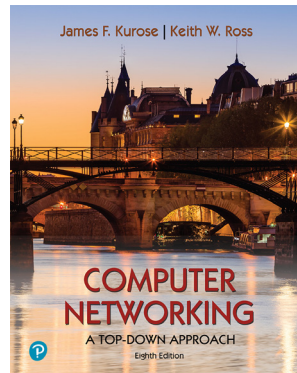
# Chapter 3

## Transport Layer

Courtesy to the textbooks' authors and Pearson Addison-Wesley because many slides are adapted from the following textbooks and their associated slides.



Jim Kurose, Keith Ross,  
“Computer Networking: A Top  
Down Approach”, 7<sup>th</sup> Edition,  
Pearson, 2016.



Jim Kurose, Keith Ross,  
“Computer Networking: A Top  
Down Approach”, 8<sup>th</sup> Edition,  
Pearson, 2020.

All material copyright 1996-2020  
J.F Kurose and K.W. Ross,  
All Rights Reserved

# Transport layer: overview

## *Our goal:*

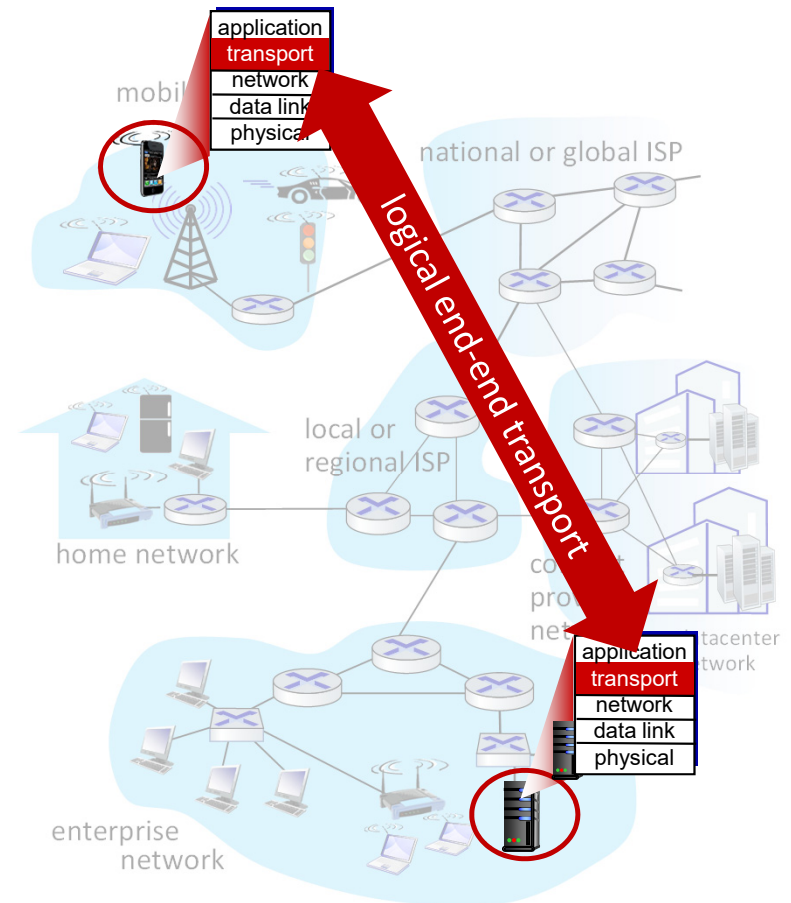
- understand principles behind transport layer services:
  - multiplexing, demultiplexing
  - reliable data transfer
  - flow control
  - congestion control
- learn about Internet transport layer protocols:
  - UDP: connectionless transport
  - TCP: connection-oriented reliable transport
  - TCP congestion control

# Transport layer: roadmap

- Review of transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control
- Evolution of transport-layer functionality

# Transport services and protocols

- provide *logical communication* between application processes running on different hosts
- transport protocols actions in end systems:
  - sender: breaks application messages into *segments*, passes to network layer
  - receiver: reassembles segments into messages, passes to application layer
- $\geq 2$  transport protocols available to Internet applications
  - TCP, UDP



# Transport vs. network layer services and protocols

- **network layer:** logical communication between *hosts*
- **transport layer:** logical communication between *processes*
  - relies on network layer services
  - enhances network layer services

## *household analogy:*

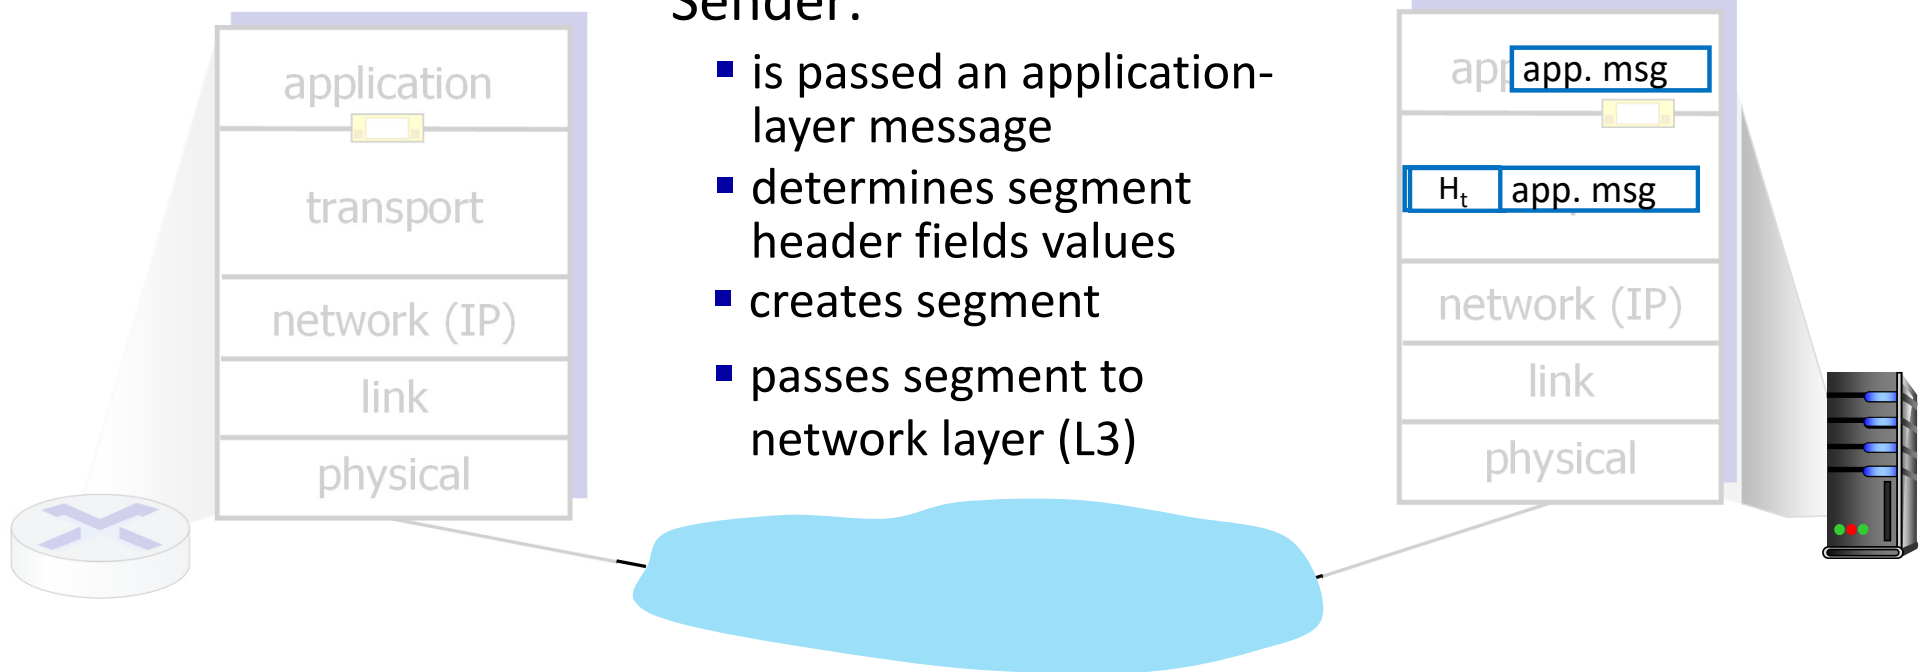
*There are multiple persons (Alice, Bob, ...) in a house:*

- hosts = houses
- processes = persons
- network-layer protocol = postal service
  - to a certain house
- transport protocol = in-site manager
  - to a certain in-house person

# Transport Layer Actions

## Sender:

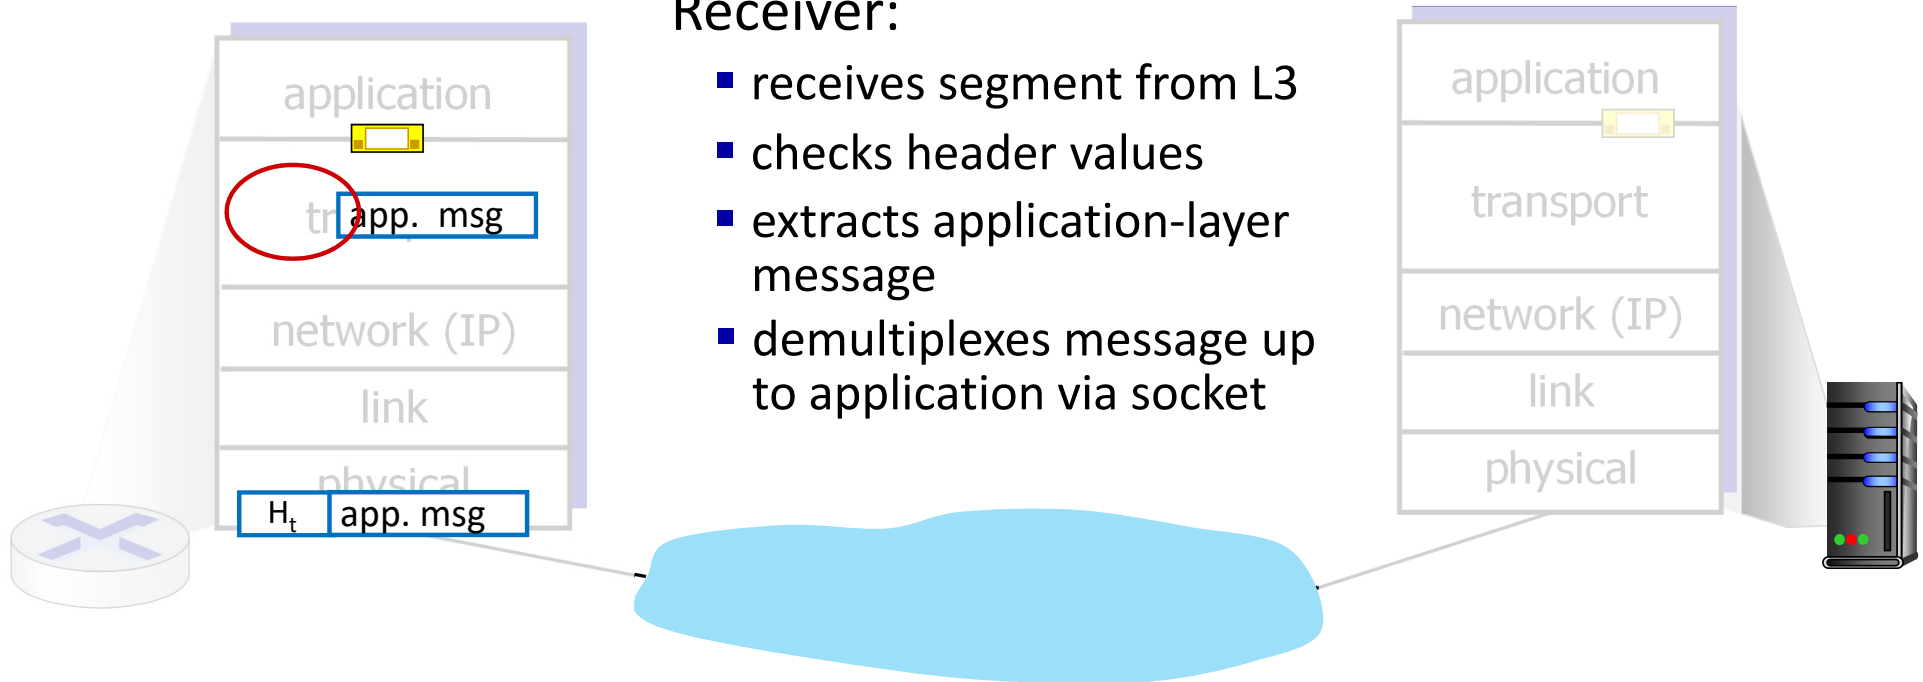
- is passed an application-layer message
- determines segment header fields values
- creates segment
- passes segment to network layer (L3)



# Transport Layer Actions

## Receiver:

- receives segment from L3
- checks header values
- extracts application-layer message
- demultiplexes message up to application via socket



# Two principal Internet transport protocols

- **TCP:** Transmission Control Protocol

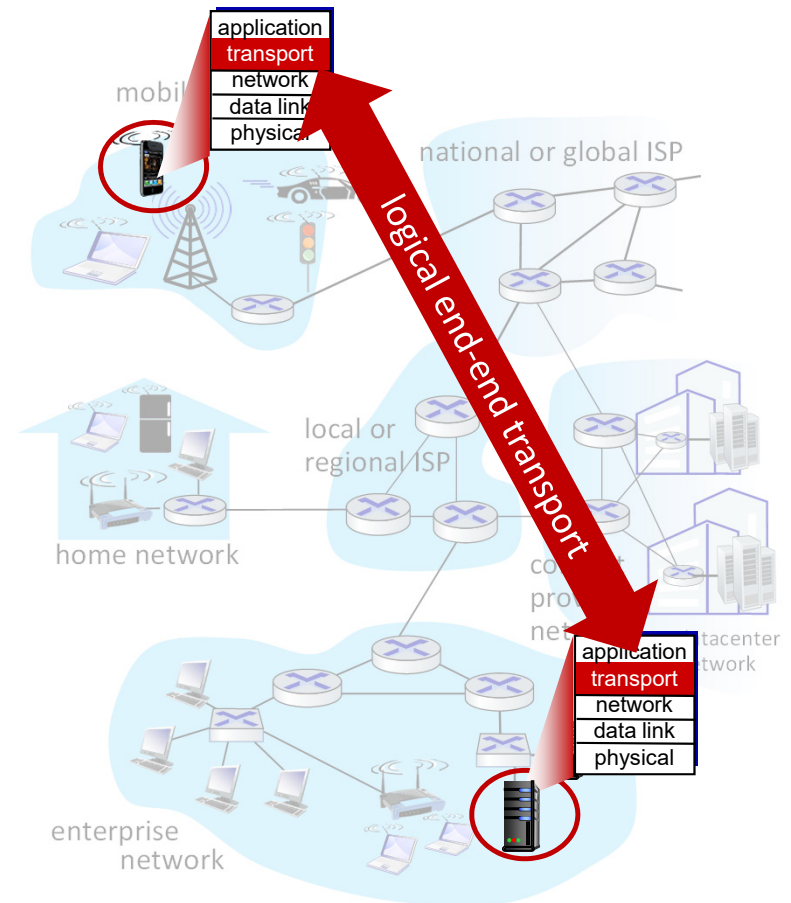
- reliable, in-order delivery
- congestion control
- flow control
- connection setup

- **UDP:** User Datagram Protocol

- unreliable, unordered delivery
- no-frills extension of “best-effort” IP

- services not available:

- delay guarantees
- bandwidth guarantees

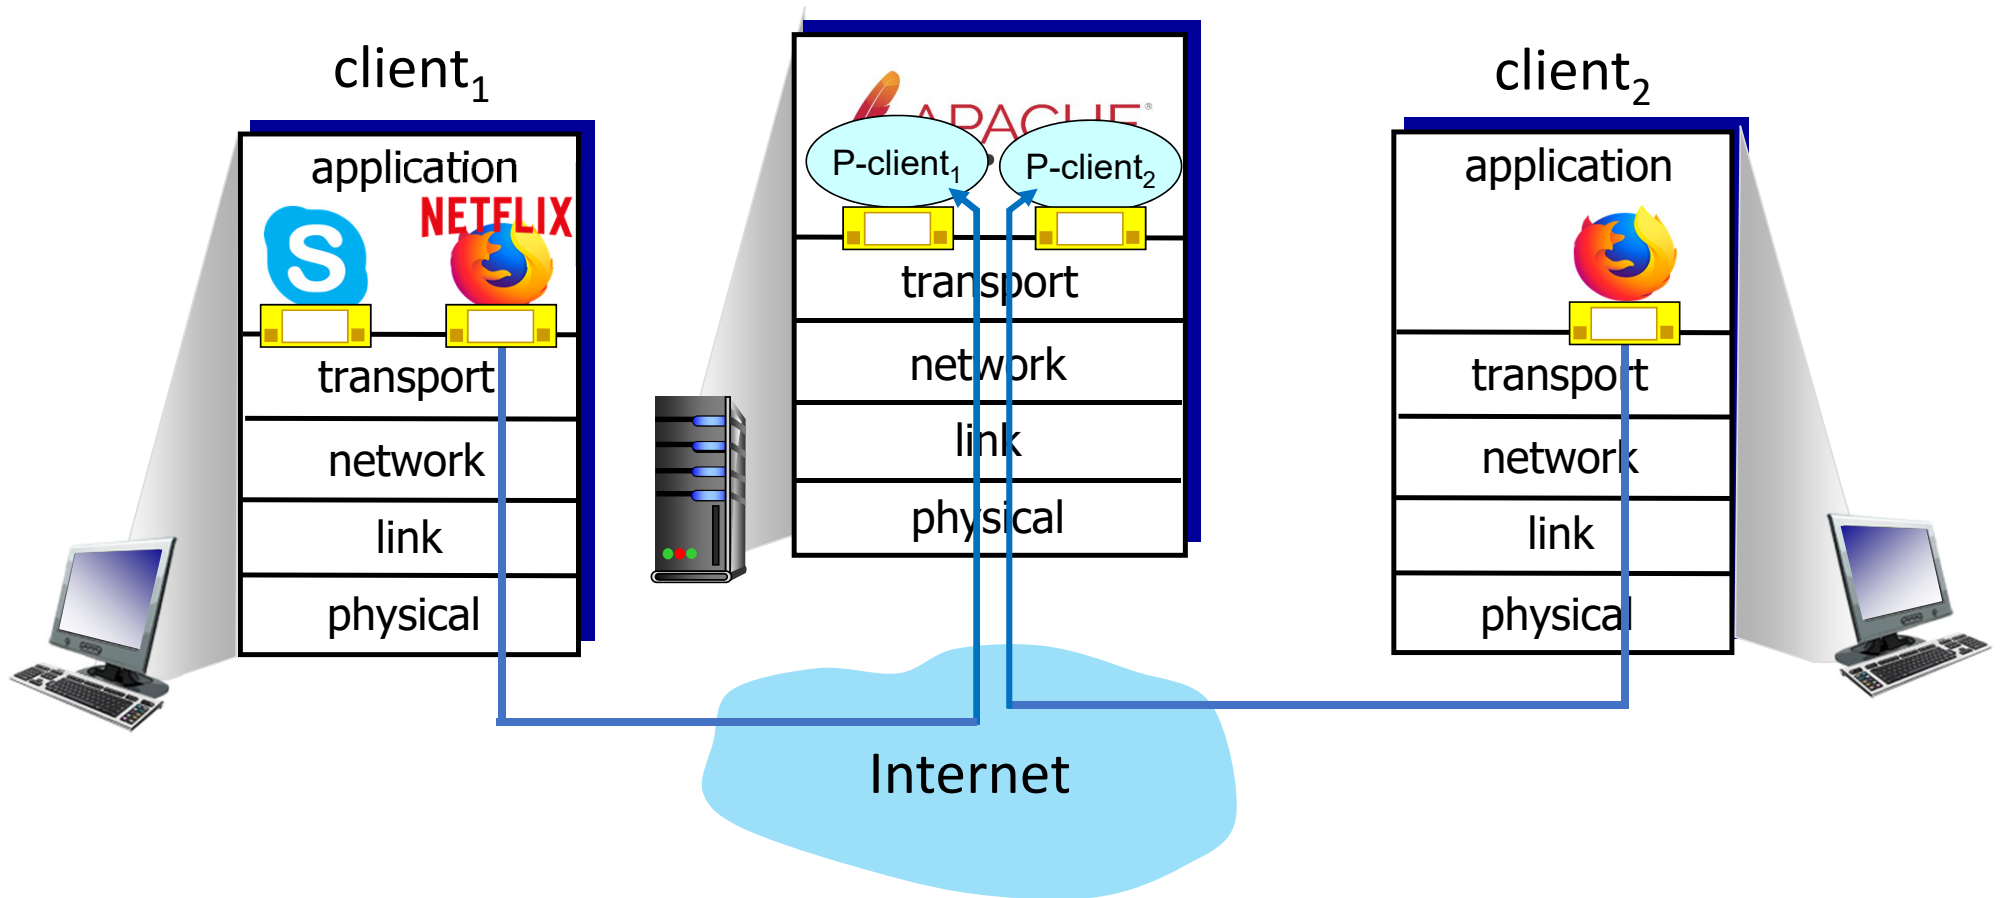




# Chapter 3: roadmap

- Review of transport-layer services
- **Multiplexing and demultiplexing**
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control
- Evolution of transport-layer functionality

## HTTP server



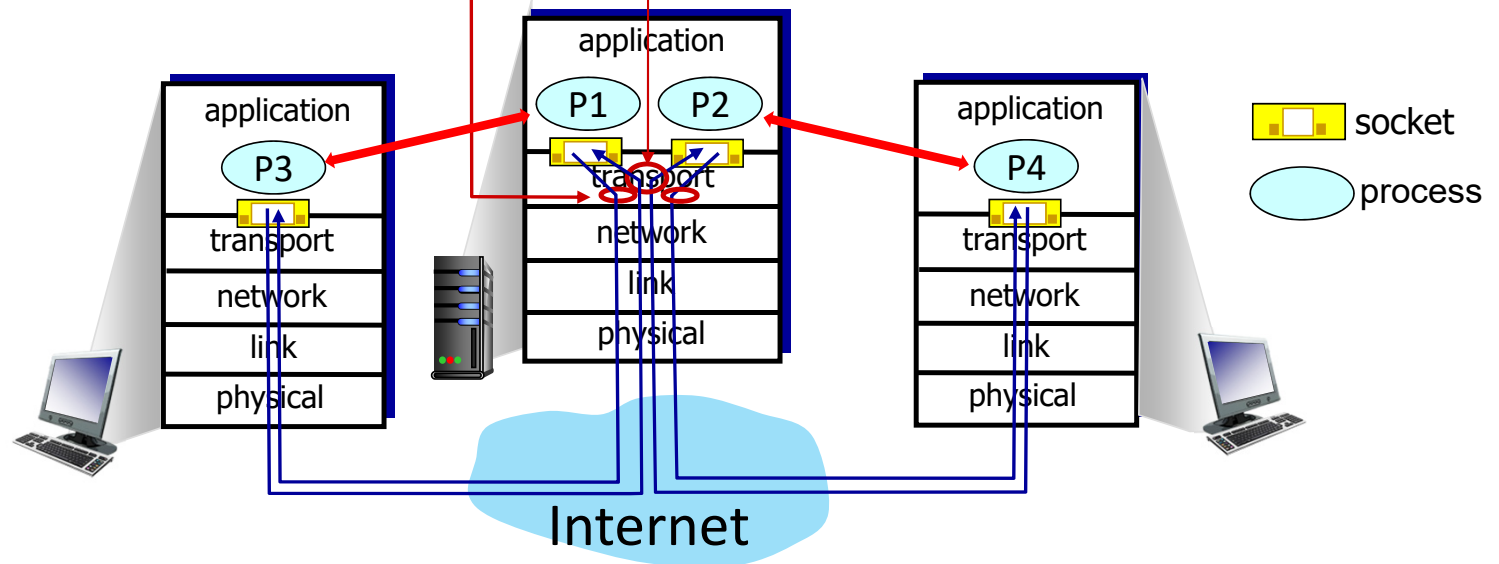
# Multiplexing/demultiplexing

## *multiplexing at sender:*

handle data from multiple sockets, add transport header (later used for demultiplexing)

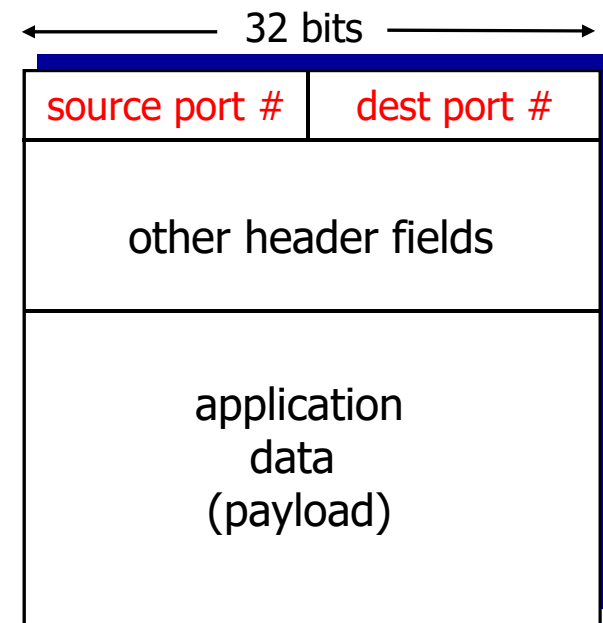
## *demultiplexing at receiver:*

use header info to deliver received segments to correct socket



# How demultiplexing works

- host receives IP datagrams
  - each datagram has source IP address, destination IP address
  - each datagram carries one transport-layer segment
  - each segment has source, destination port number
- host uses *IP addresses* & *port numbers* to direct segment to appropriate socket/process



TCP/UDP segment format

# Connectionless (UDP) demultiplexing

*Recall:*

- when creating socket, must specify *host-local* port #:

```
DatagramSocket mySocket1  
= new DatagramSocket(12534);
```

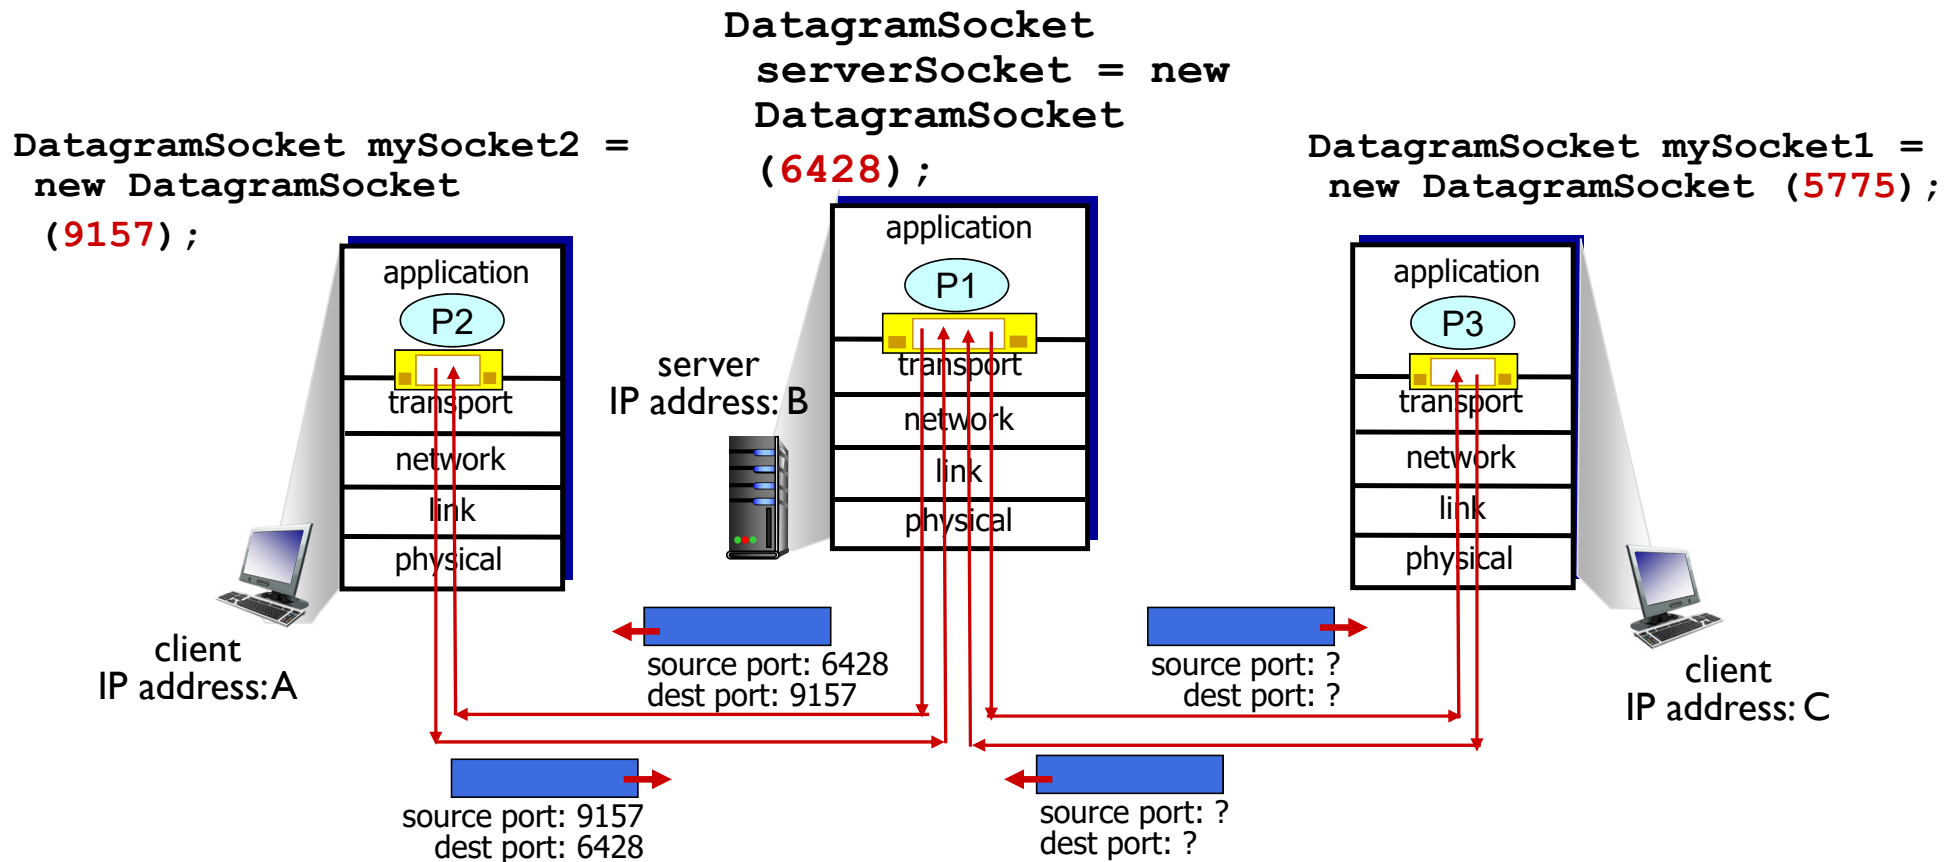
- when app at the source side creates datagram to send into UDP socket, it must specify
  - destination IP address
  - destination port #

- when receiving host receives UDP segment:
  - checks destination port # in segment
  - directs UDP segment to socket with that port #



UDP datagrams with *same dest. port #*, but different source IP addresses and/or source port numbers will be directed to *same socket* at receiving host

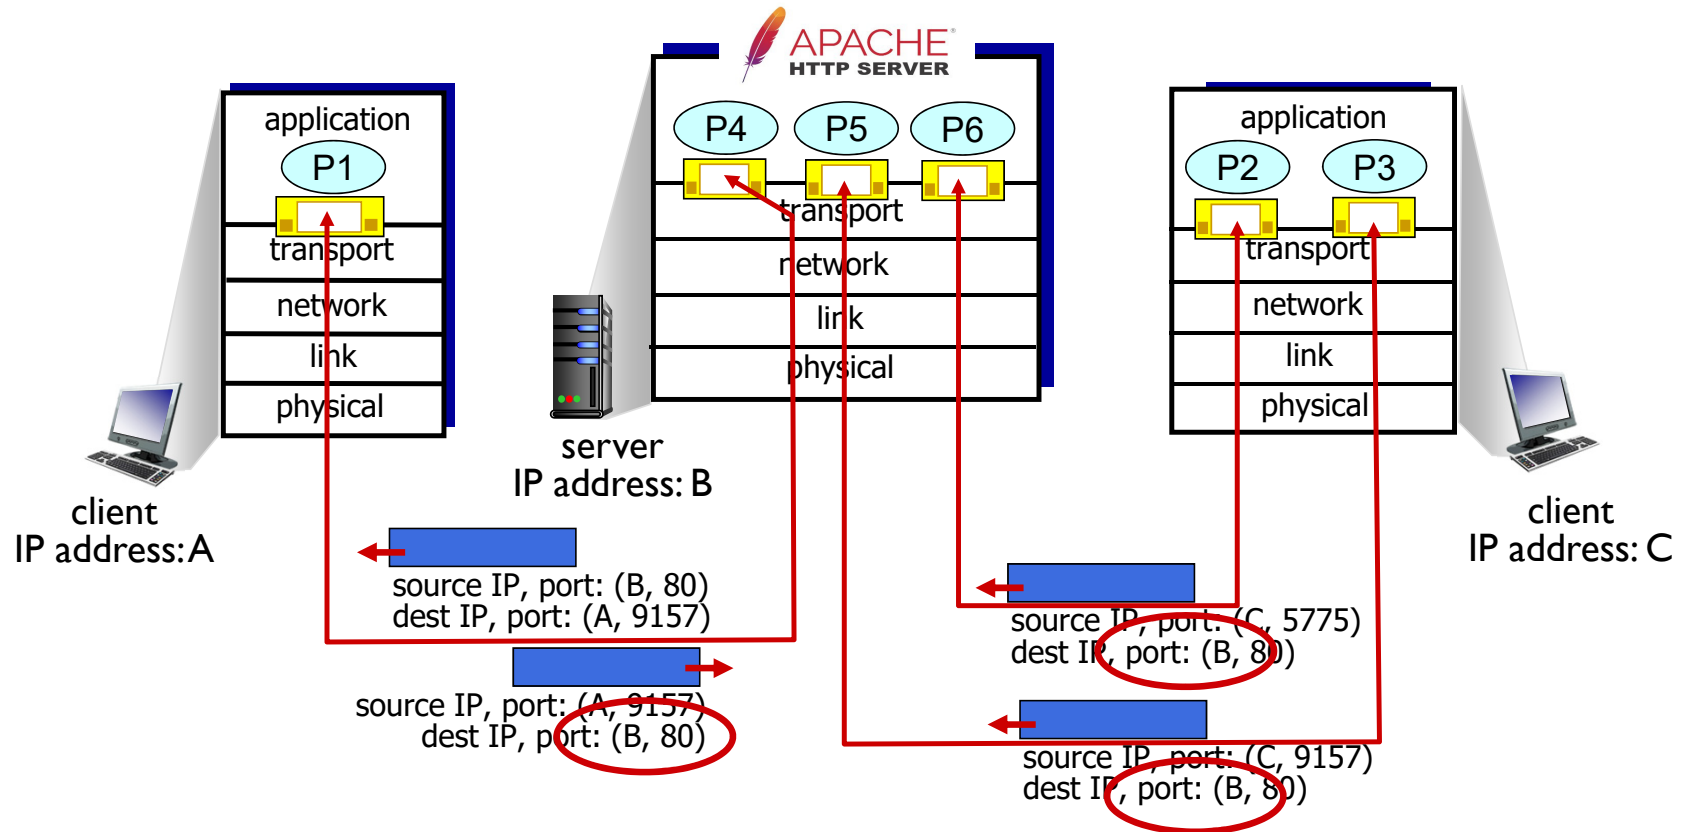
# Connectionless demultiplexing: an example



# Connection-oriented (TCP) demultiplexing

- TCP socket identified by 4-tuple:
  - source IP address
  - source port number
  - dest IP address
  - dest port number
- demux: receiver uses *all four values (4-tuple)* to direct segment to appropriate socket
- server may support many simultaneous TCP sockets:
  - each socket identified by its own 4-tuple
  - each socket associated with a different connecting client

# Connection-oriented demultiplexing: example



Three segments, all destined to (B,80) are demultiplexed to *different* sockets



# Summary

- Multiplexing and demultiplexing are based on header field values
- **UDP:** destination host demultiplexes segments using destination port number (only)
  - a UDP socket is uniquely identified by (dest IP, dest port #)
- **TCP:** demultiplexing using 4-tuple
  - IP addresses and port numbers of source and destination
- Multiplexing/demultiplexing happen at *all* layers

# Chapter 3: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- **Connectionless transport: UDP**
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control
- Evolution of transport-layer functionality



# UDP: User Datagram Protocol

- “no frills,” “bare bones” Internet transport protocol
- “best effort” service, UDP segments may be:
  - lost
  - delivered out-of-order to app
- *connectionless*:
  - no handshaking between (UDP) sender and receiver
  - each UDP segment is handled independently of others

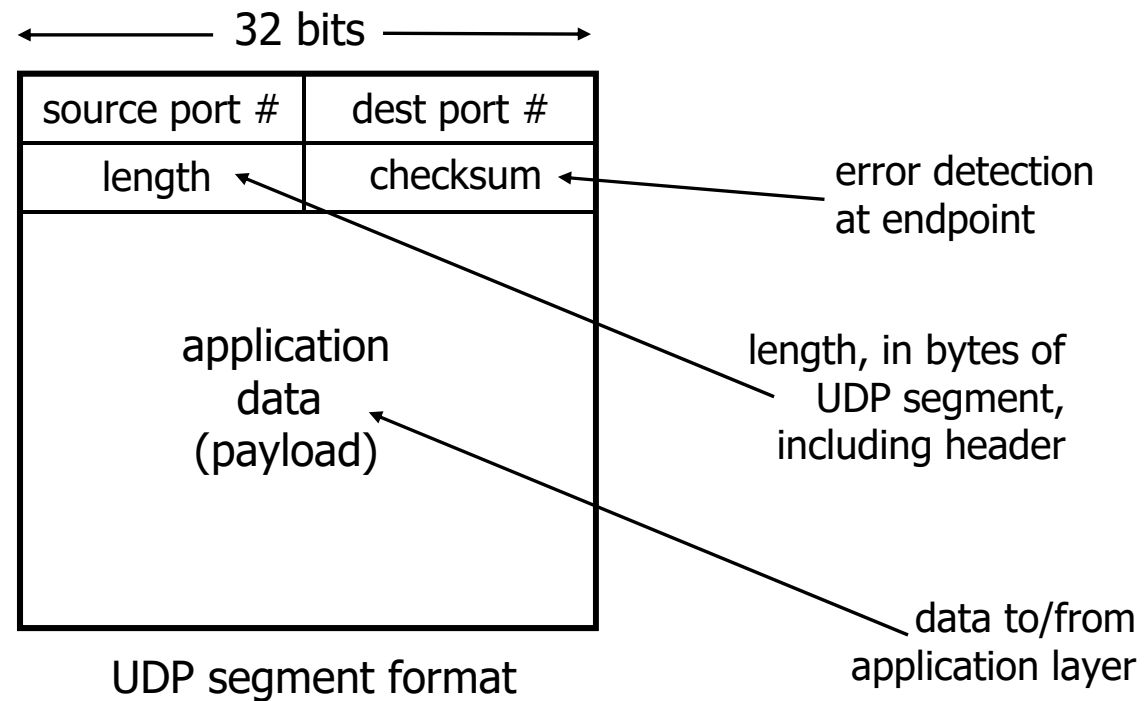
## Why is there a UDP?

- no connection establishment
  - which can add RTT delay
- simple (no connection state at endpoint)
  - server can support more active clients when running over UDP
- small header overhead
- finer application-layer control
  - no congestion control
  - UDP can blast away as fast as desired!

# UDP: User Datagram Protocol

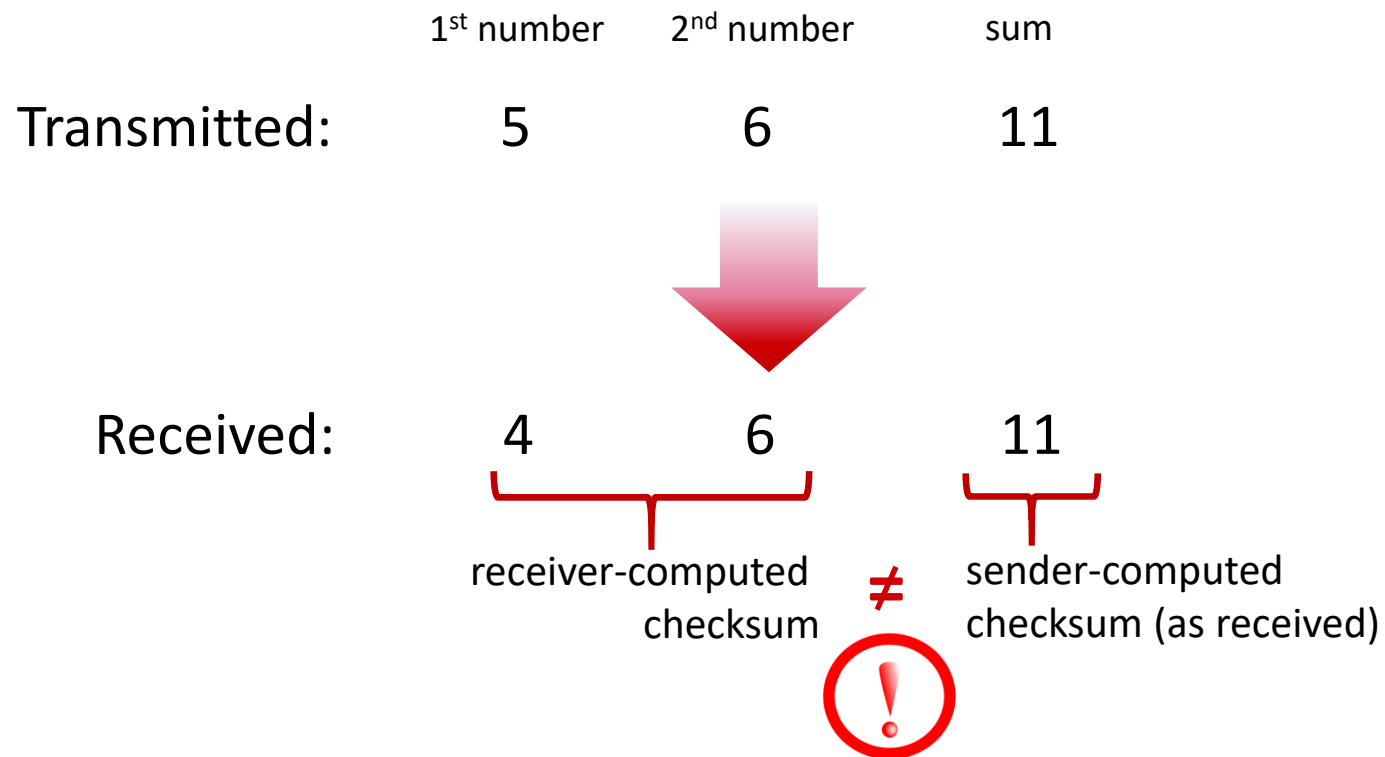
- UDP was ratified in 1980. [RFC 768]
- UDP is used by:
  - streaming multimedia apps (loss tolerant, rate sensitive)
  - DNS
  - SNMP
  - HTTP/3 (over QUIC over UDP)
- if reliable transfer needed over UDP (e.g., HTTP/3):
  - add needed reliability at application layer
  - add congestion control at application layer

# UDP segment format



# UDP checksum

*Goal:* detect errors (*i.e.*, flipped bits) in transmitted segment



# UDP checksum

*Goal:* detect errors (*i.e.*, flipped bits) in transmitted segment

## sender:

- treat contents of UDP segment as sequence of 16-bit integers
  - including pseudo header (IP addresses, ...)
- **checksum:**
  - add segment content (1's complement sum)
  - and then take 1's complement
- checksum value put into UDP checksum field

## receiver:

- compute checksum of received segment
- check whether computed checksum equals checksum field value:
  - Not equal – error detected
  - Equal – no error detected. *But maybe errors nonetheless?* More later ....

# Internet checksum: an example

example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
	<hr/>															
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
	<hr/>															
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

*Note:* when adding numbers, a carryout from the most significant bit needs to be added to the result

\* Check out the online interactive exercises for more examples: [http://gaia.cs.umass.edu/kurose\\_ross/interactive/](http://gaia.cs.umass.edu/kurose_ross/interactive/)



# Internet checksum: weak protection!

example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
	<hr/>															
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

Even though numbers have changed (bit flips), *no* change in checksum!

# Summary: UDP

- “no frills” protocol:
  - segments may be lost, delivered out of order
  - best effort service: “send and hope for the best”
- UDP has its pluses:
  - no setup/handshaking needed (no RTT incurred)
  - smaller header overhead
  - finer application-layer control
  - helps with reliability (checksum)
- If needed, additional functionality can be built on top of UDP in application layer (e.g., HTTP/3)