


# Go-Back-N: sender side

- **cumulative ACK:** ACK( $n$ ): ACKs all packets up to, including seq #  $n$
- sender: “sliding window” of up to  $N$  (consecutive sent but unACKed) pkts
  - k-bit seq # in pkt header

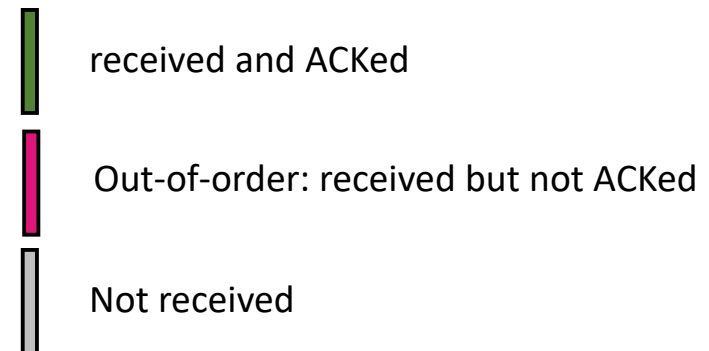



- on receiving ACK( $n$ ): move window forward to begin at  $n+1$
- timer for oldest unACKed packet.
  - timer starts if either (non-empty) window moves forward or app-layer sends a pkt to empty window
- *timeout( $n$ )*: retransmit packet  $n$  and all (available) packets in window

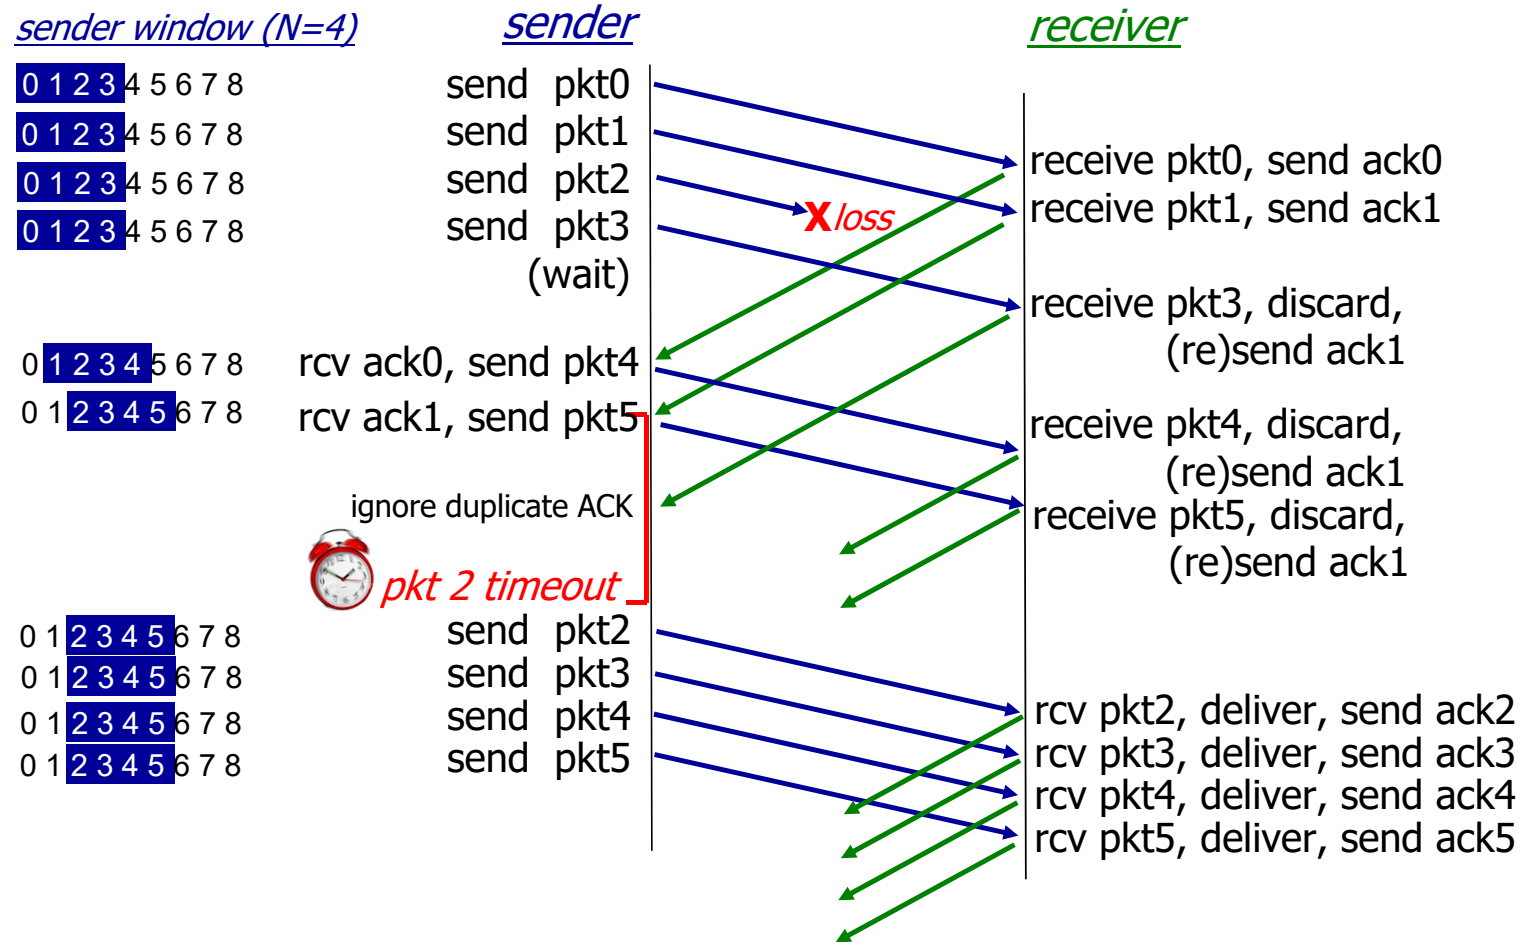
# Go-Back-N: receiver side

- Always send ACK for correctly-received packet, with highest *in-order* seq #
  - need only remember `rcv_base`
  - might generate duplicate ACKs
- on receipt of in-order packet:
  - update `rcv_base`
  - send ACK with the (in-order) seq #
- on receipt of out-of-order packet:
  - discard (don't buffer) or buffer the packet, which is an implementation decision
  - re-ACK with highest in-order seq #

Receiver view of sequence number space:



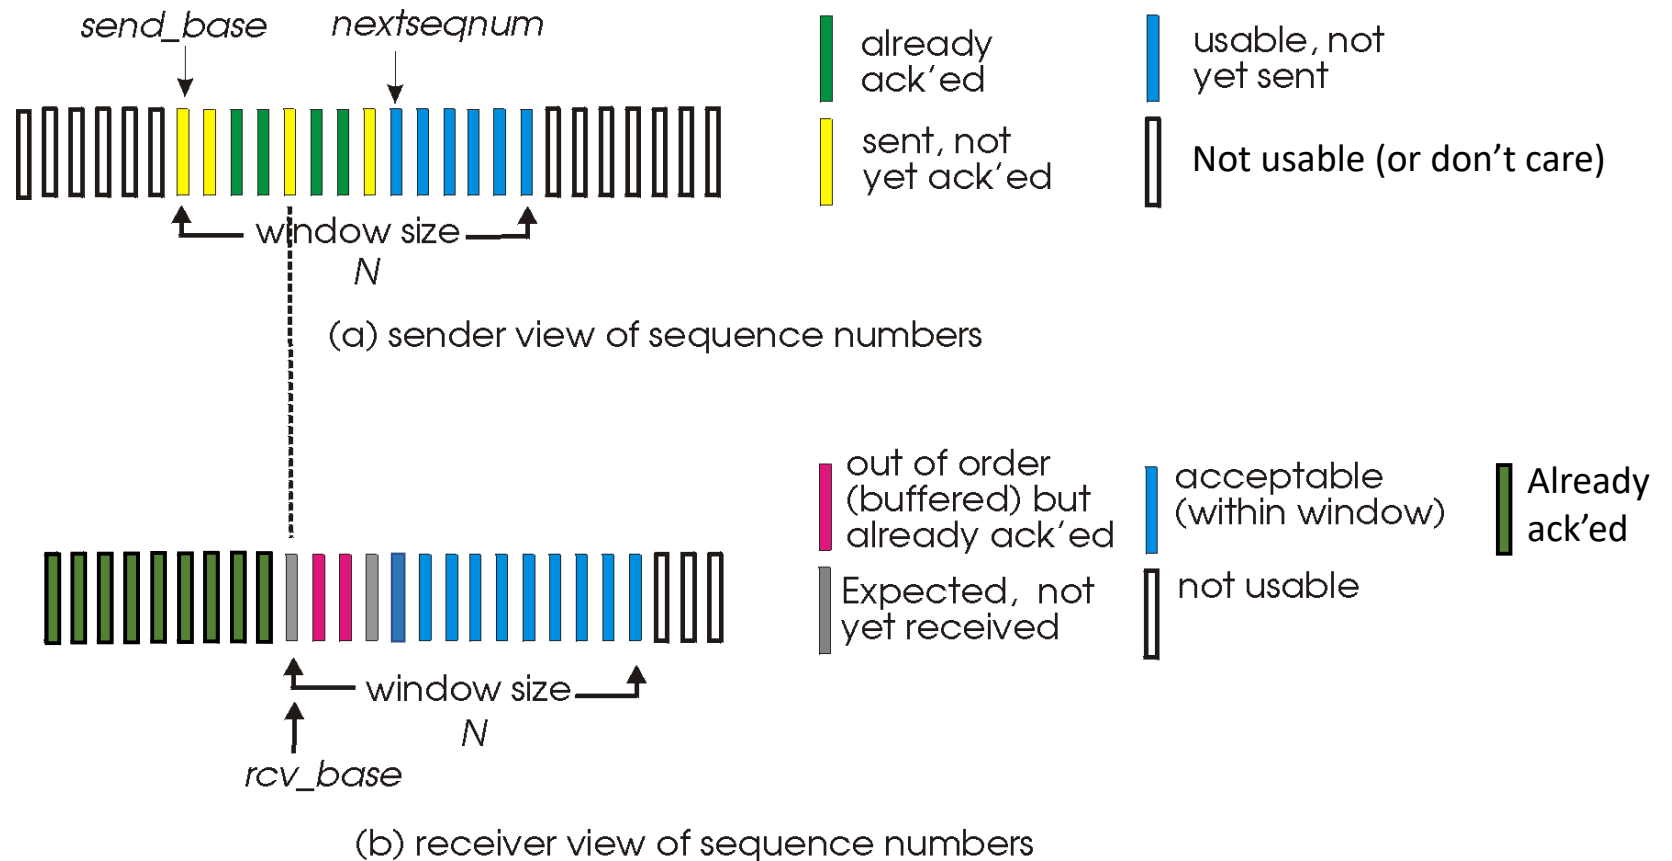
# Go-Back-N in action



# Selective repeat

- receiver *individually* acknowledges all correctly received packets
  - buffers packets, as needed, for eventual in-order delivery to upper layer
- sender times-out/retransmits individually for unACKed packets
  - sender maintains timer for each unACKed pkt
- sender window
  - $N$  consecutive seq #
  - limits # of sent and still unACKed packets

# Selective repeat: sender, receiver windows



# Selective repeat: sender and receiver

## sender

### data from above:

- if next available seq # in window, send packet

### timeout( $n$ ):

- resend packet  $n$ , restart timer

### ACK( $n$ ) in

[sendbase, sendbase+N-1]:

- mark packet  $n$  as received
- if  $n$  is oldest unACKed packet, advance window base to next unACKed seq #

## receiver

### packet $n$ in [rcvbase, rcvbase+N-1]

- send ACK( $n$ )
- out-of-order: buffer
- in-order: deliver (also deliver buffered, in-order packets), advance window to next not-yet-received packet

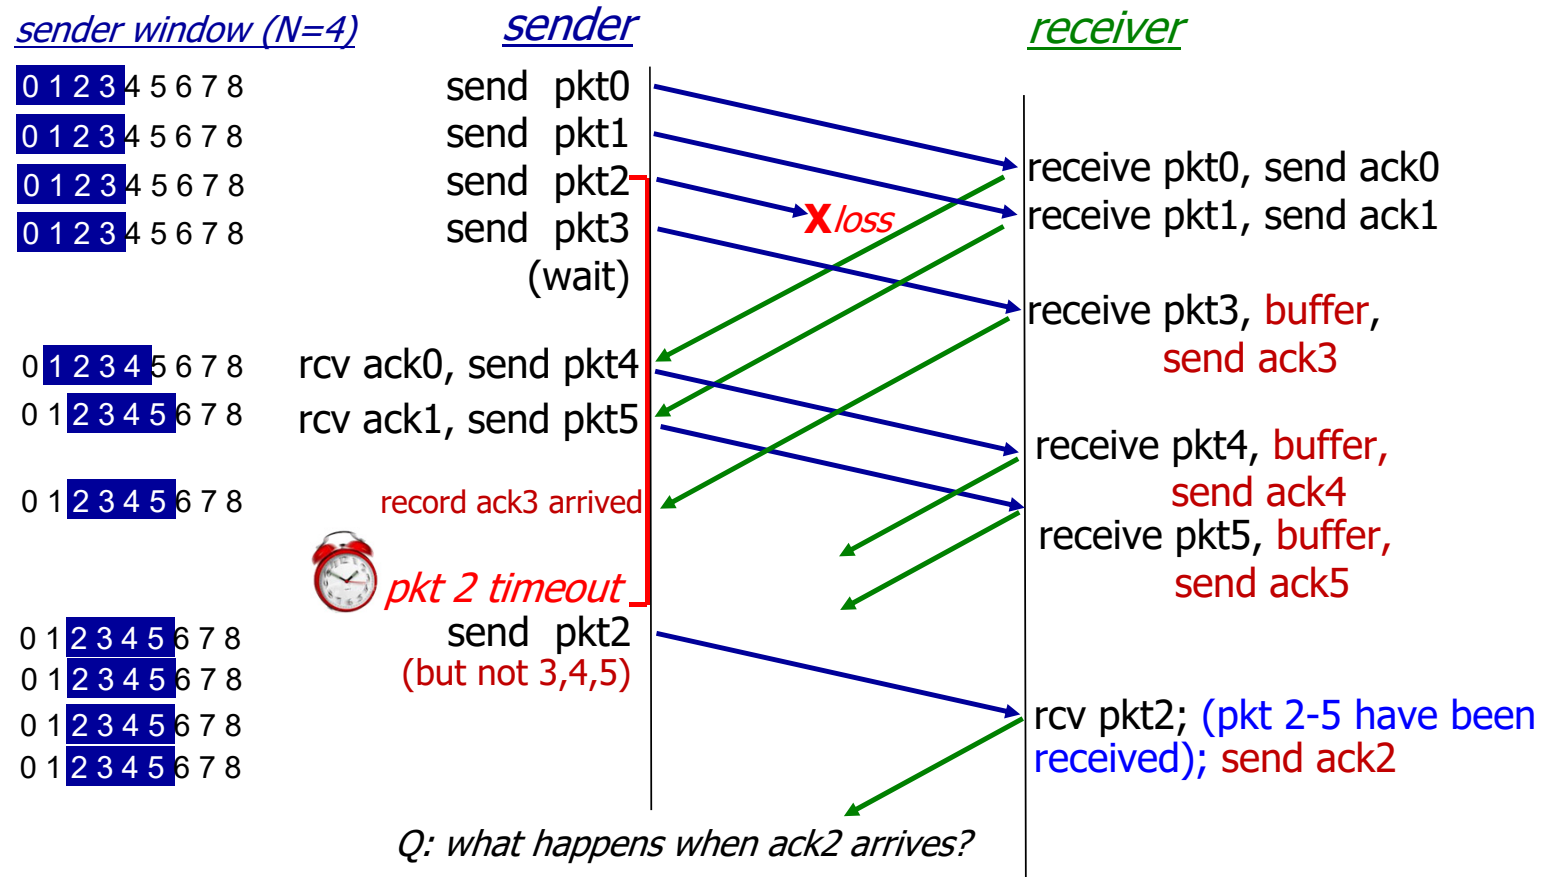
### packet $n$ in [rcvbase-N, rcvbase-1]

- send ACK( $n$ ) back to the sender

### otherwise:

- ignore

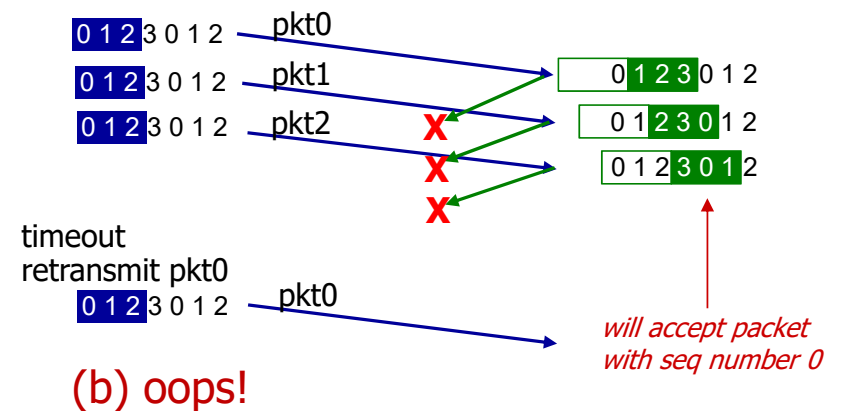
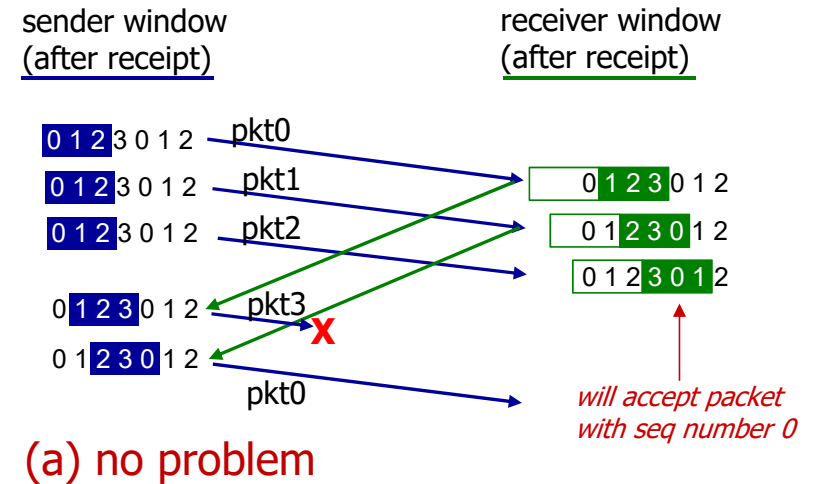
# Selective Repeat in action



# Selective repeat: a dilemma!

example:

- use 2-bit seq # : {0, 1, 2, 3}
- window size = 3



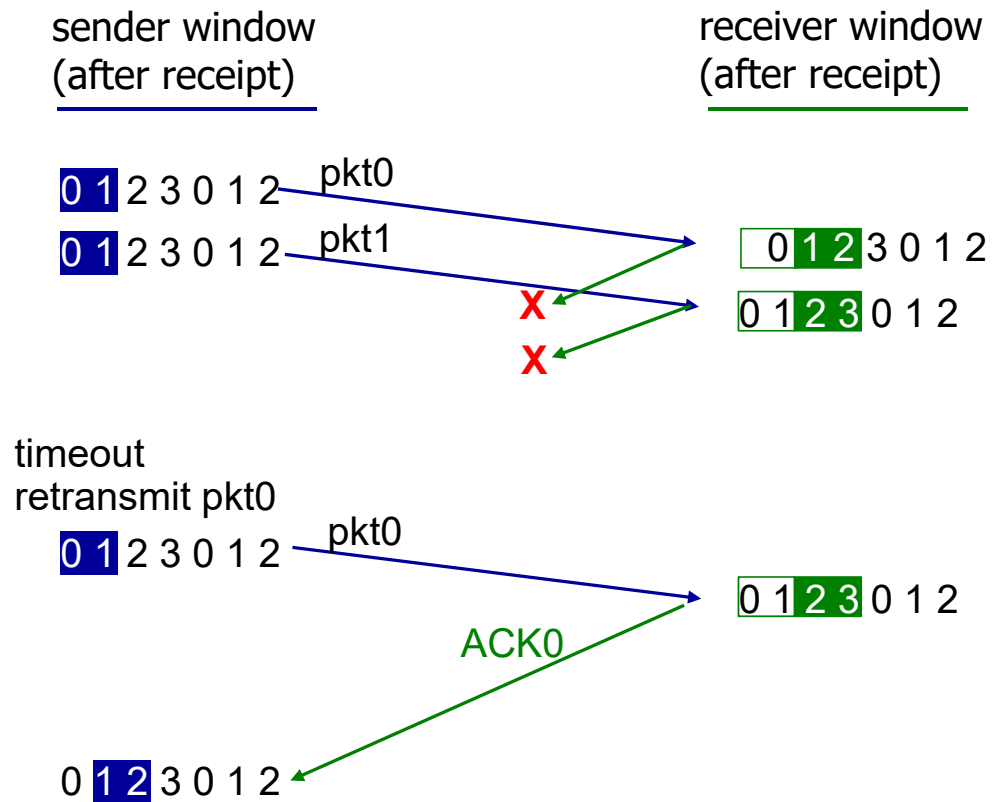


# Selective repeat: a dilemma!

example:

- use 2-bit seq # : {0, 1, 2, 3}
- window size = 2

**Q:** what relationship is needed between sequence # size and window size to avoid the problem in scenario (b)?



# Chapter 3: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- **Connection-oriented transport: TCP**
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- Principles of congestion control
- TCP congestion control

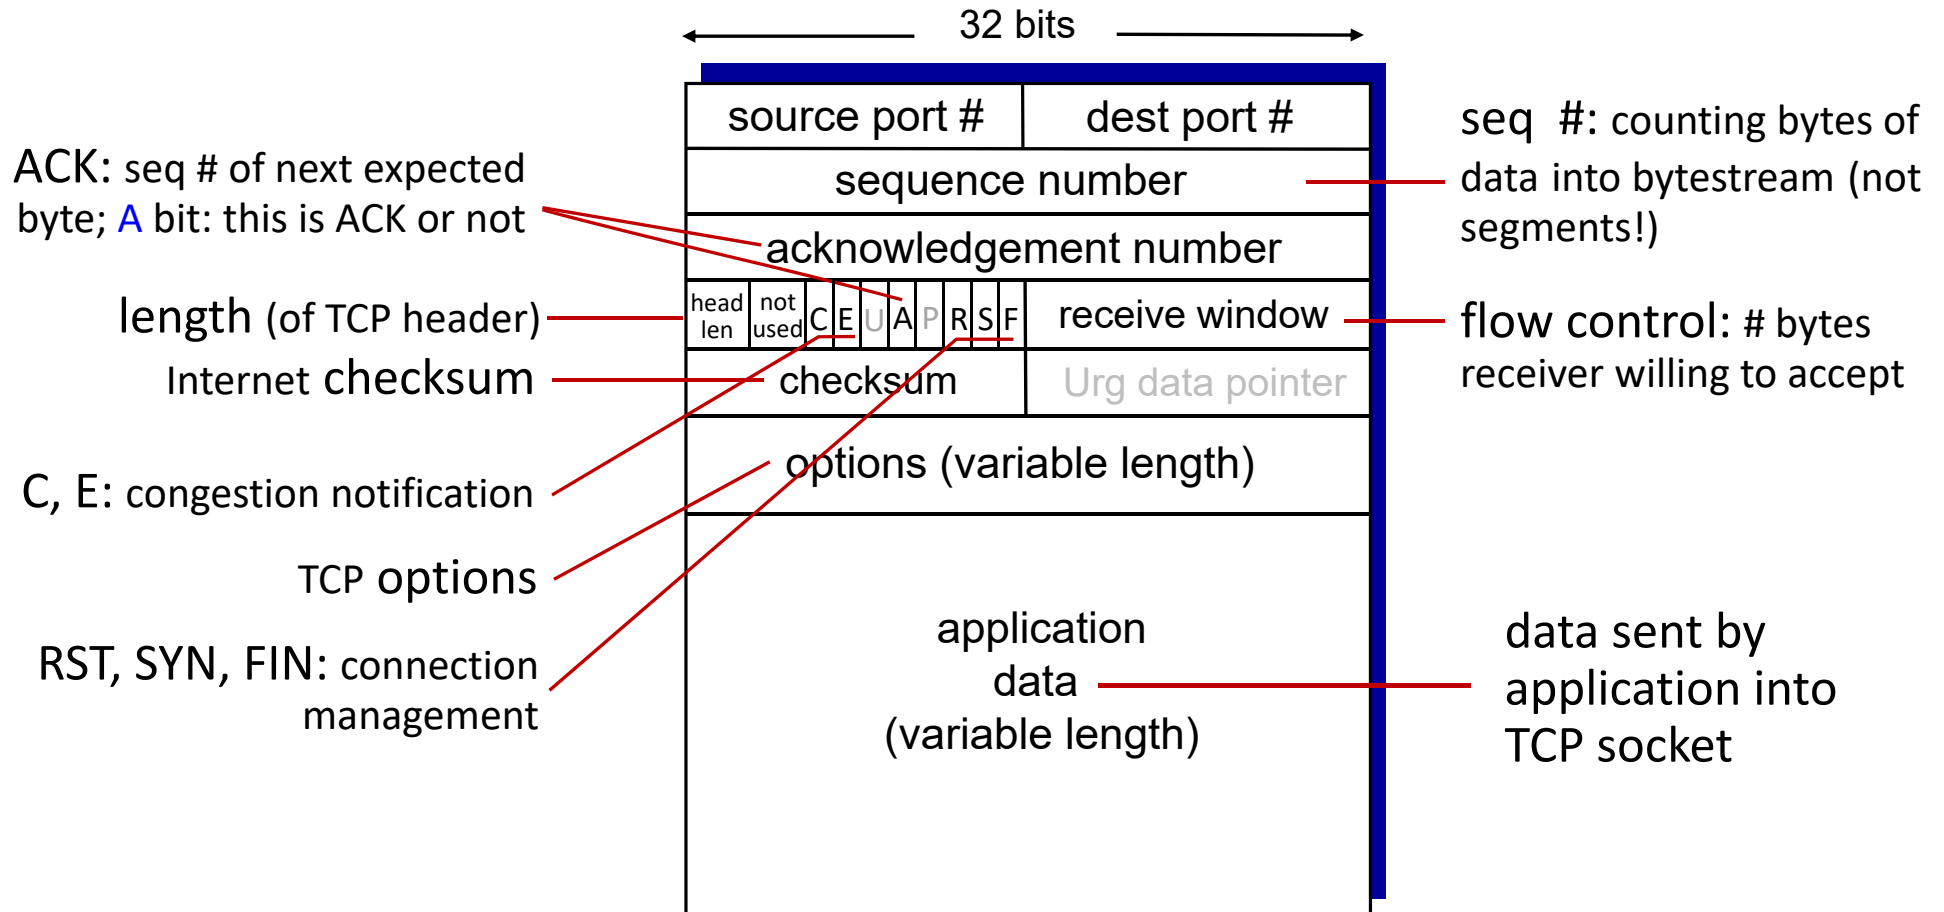


# TCP: overview

RFCs: 793, 1122, 2018, 5681, 7323

- **point-to-point:**
  - one sender, one receiver
- **reliable, in-order *byte stream*:**
  - no “message boundaries”
  - MSS: maximum segment size
- **full duplex data:**
  - bi-directional data flow in same (TCP) connection
- **cumulative ACKs**
- **pipelining:**
  - TCP congestion and flow control set window size
- **connection-oriented:**
  - handshaking (exchange of control messages) initializes sender, receiver state before data exchange
- **flow controlled:**
  - sender will not overwhelm receiver

# TCP segment structure



# TCP sequence numbers and ACKs

## Sequence numbers:

- byte stream “number” of first byte in segment’s data

## Acknowledgements:


- seq # of next byte expected from other side
- cumulative ACK

Q: how receiver handles out-of-order segments

- A: TCP spec doesn’t say, - up to implementor

outgoing segment from sender

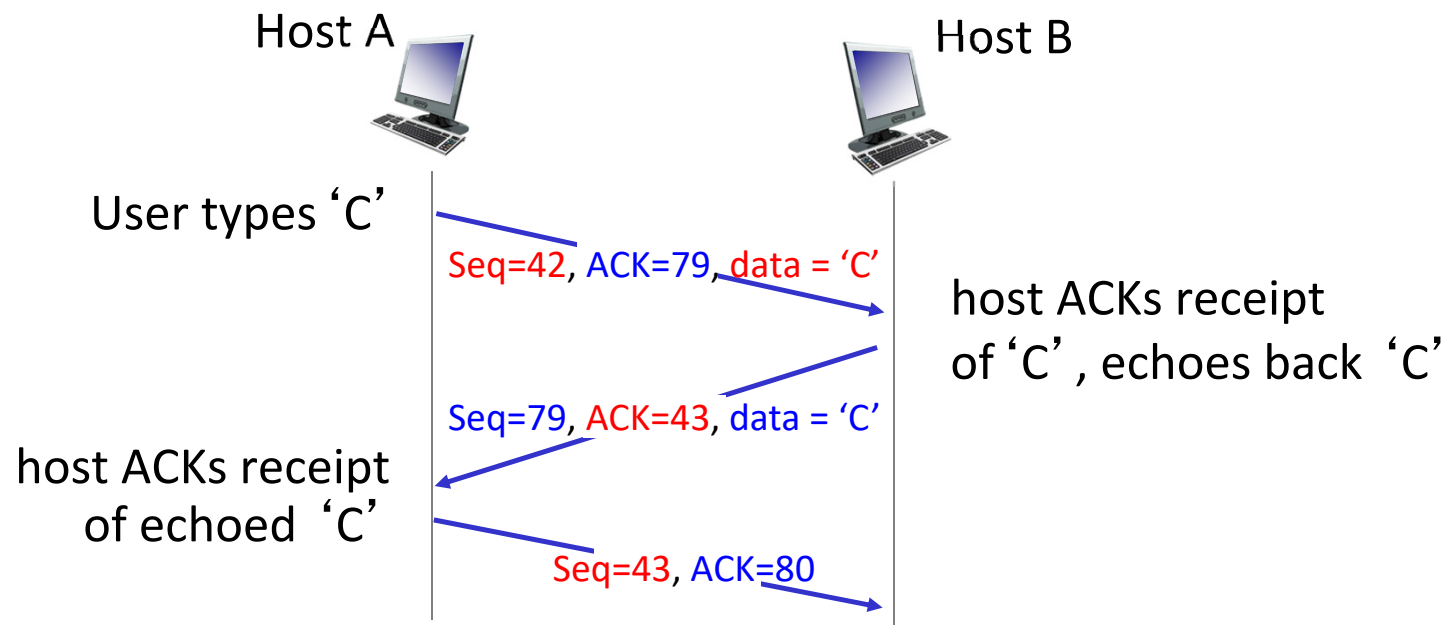
source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer



outgoing segment from receiver

source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer

# TCP sequence numbers and ACKs



simple telnet scenario

# TCP round trip time and timeout

Q: how to set TCP timeout value?

- longer than RTT, but RTT varies!
- *too short*: premature timeout, unnecessary retransmissions
- *too long*: slow reaction to segment loss


Q: how to estimate RTT?

- `SampleRTT`: measured time from segment transmission until ACK receipt
  - ignore retransmissions
- `SampleRTT` will vary and we want estimated RTT “smoother”
  - average several *recent* measurements, not just current `SampleRTT`

# TCP round trip time and timeout

$$\text{EstimatedRTT} := (1-\alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- exponential weighted moving average (EWMA)
- influence of past sample decreases exponentially fast
- typical value:  $\alpha = 0.125$





# TCP round trip time and timeout

- timeout interval: **EstimatedRTT** plus “safety margin”
  - large variation in **EstimatedRTT**: want a larger safety margin

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$



↑  
estimated RTT

↑  
“safety margin”

- **DevRTT**: EWMA of **SampleRTT** deviation from **EstimatedRTT**:

$$\text{DevRTT} := (1 - \beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically,  $\beta = 0.25$ )

\* Check out the online interactive exercises for more examples: [http://gaia.cs.umass.edu/kurose\\_ross/interactive/](http://gaia.cs.umass.edu/kurose_ross/interactive/)

# TCP Sender (simplified)

event: data received from application layer

- create segment with seq #
- seq # is byte-stream number of first data byte in segment
- start timer if not already running
  - think of timer as for oldest unACKed segment
  - expiration interval: **TimeoutInterval**

event: timeout

- retransmit the (only one) segment that caused timeout
- restart timer

event: ACK received


- if ACK acknowledges previously unACKed segments
  - update what is known to be ACKed (window moves forward)
  - start timer if there are still unACKed segments

# TCP Receiver: ACK generation [RFC 5681]


<i>Event at receiver</i>	<i>TCP receiver action</i>
arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK
arrival of in-order segment with expected seq #. One other segment has ACK pending	immediately send single cumulative ACK, ACKing both in-order segments
arrival of out-of-order segment higher-than-expect seq. # . Gap detected	immediately send <i>duplicate ACK</i> , indicating seq. # of next expected byte
arrival of segment that partially or completely fills gap	immediate send ACK, provided that segment starts at lower end of gap

# TCP: retransmission scenarios

lost ACK scenario

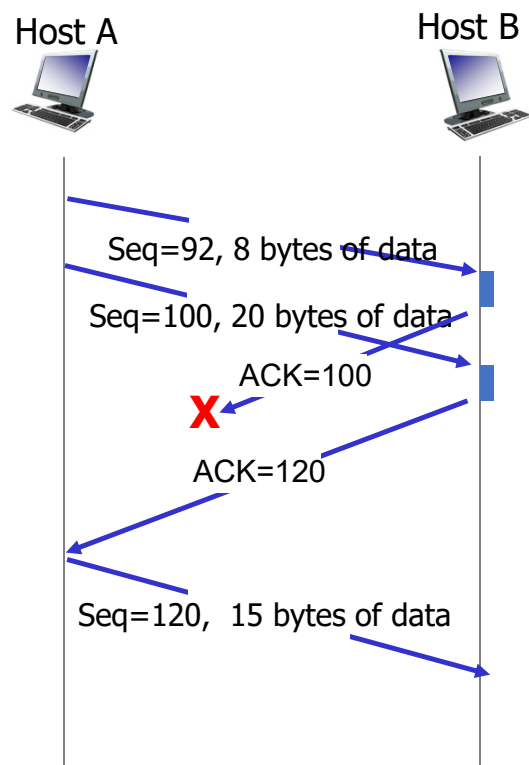


premature timeout



# TCP: retransmission scenarios

cumulative ACK covers  
for earlier lost ACK



# TCP fast retransmit

## *TCP fast retransmit*

if sender receives 3 additional ACKs for same data (“triple duplicate ACKs”), resend unACKed segment with smallest seq #

- likely that unACKed segment lost, so don't wait for timeout

💡 Receipt of three duplicate ACKs indicates 3 segments received after a missing segment – lost segment is likely. So retransmit!

