



# Chapter 6

## The Link Layer and LANs

Courtesy to the textbooks' authors and Pearson Addison-Wesley because many slides are adapted from the following textbooks and their associated slides.



Jim Kurose, Keith Ross,  
“Computer Networking: A Top  
Down Approach”, 7<sup>th</sup> Edition,  
Pearson, 2016.



Jim Kurose, Keith Ross,  
“Computer Networking: A Top  
Down Approach”, 8<sup>th</sup> Edition,  
Pearson, 2020.

All material copyright 1996-2020  
J.F Kurose and K.W. Ross,  
All Rights Reserved


# Link layer, LANs: roadmap

- introduction
- error detection, correction
- multiple access protocols
- LANs
  - addressing, ARP
  - Ethernet
  - switches
  - VLANs
- link virtualization: MPLS
- data center networking
- a day in the life of a web request




# Link layer: introduction

- Terminology
  - host and router are  $\geq$ L3 nodes
    - switch (and bridge) is layer-2
  - frame is layer-2 packet
- **link layer** is responsible to transfer packets from one node to another via a **link** (within a **subnet**)
  - either wired or wireless link
  - LAN



# Link layer: context

- datagram transferred by different link-layer protocols over different links:
  - e.g., WiFi on first link, Ethernet on next link



- each link protocol provides different services
  - e.g., may or may not provide reliable data transfer over link

# Link layer: services

## ■ framing

- encapsulate datagram into frame
  - adding header, trailer
- “MAC” addresses in frame headers identify source, destination
  - MAC address is different from IP address!




## ■ channel access (multiple access, MAC)

- if shared medium

## ■ reliable delivery in link layer

- we already know how to do this!
- seldom used on low bit-error links
- wireless links: high error rates
- Q: why both link-level and end-to-end reliability?

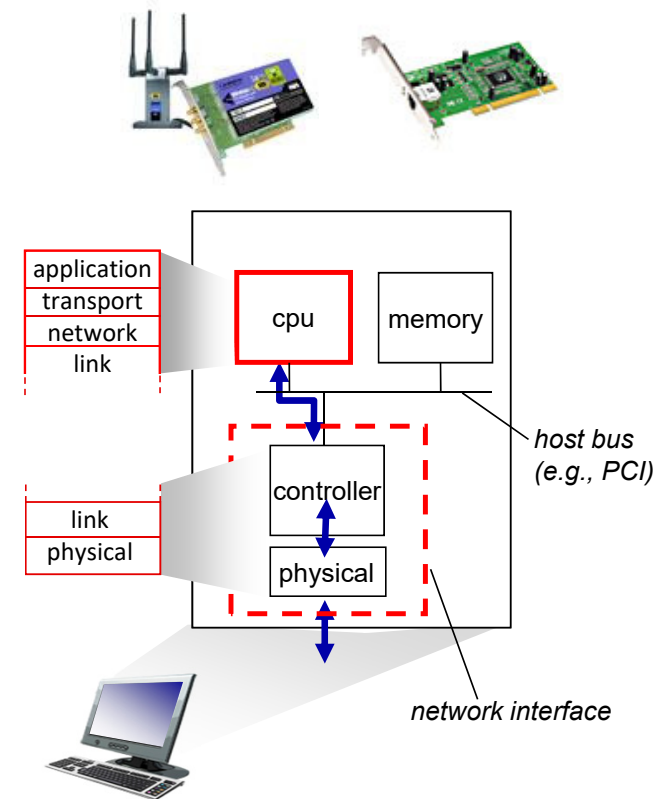


## Link layer: services (more)


- **flow control:**
  - pacing between adjacent sending and receiving nodes
- **error detection:**
  - errors caused by signal attenuation, noise, collision.
  - retransmits or drops frame, if receiver detects errors
- **error correction:**
  - receiver identifies *and corrects* bit error(s) without retransmission
- **half-duplex and full-duplex:**
  - with half duplex, nodes at both ends of link can transmit, but not at same time

# Where is the link layer implemented?

- in each host
- link layer implemented in *network interface card* (NIC)
  - (Ethernet, WiFi) card or chip
  - implements link+physical layers
- attaches into host's system bus
- combination of hardware, software, firmware



# Interfaces communicating



sending side:

- encapsulates datagram in frame
- adds error checking bits, reliable data transfer, flow control, etc.

receiving side:

- looks for errors, reliable data transfer, flow control, etc.
- extracts datagram, passes to upper layer at receiving side



# Link layer, LANs: roadmap

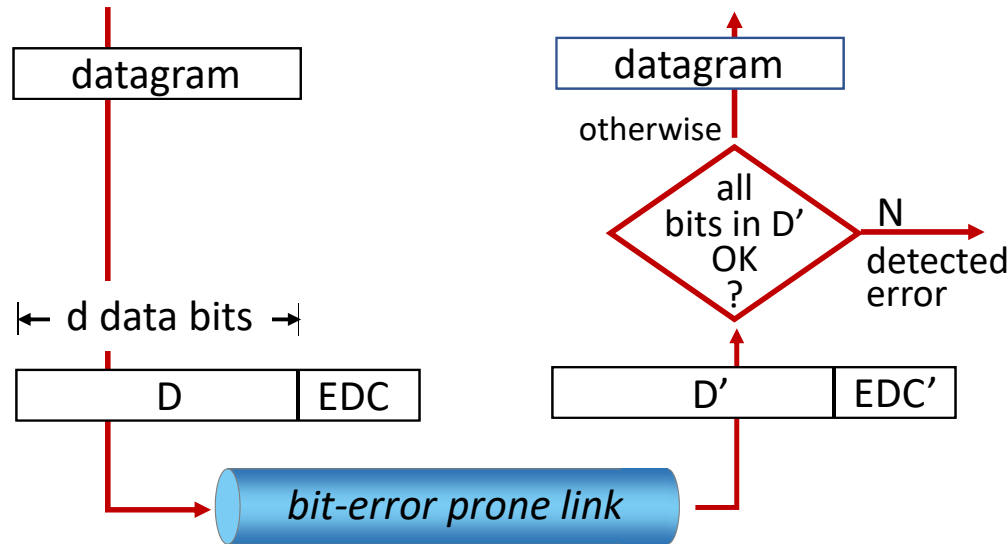
- introduction
- error detection, correction
- multiple access protocols
- LANs
  - addressing, ARP
  - Ethernet
  - switches
  - VLANs
- link virtualization: MPLS
- data center networking
- a day in the life of a web request



# Error detection

EDC: error detection and correction bits (e.g., redundancy)

D: data protected by error checking, may include header fields




Error detection has its capability

- protocol may miss some errors, but rarely
- larger EDC field yields better detection
  - same as correction

# Parity checking

## single bit parity:

- detect single bit errors




**Even parity:** set parity bit so there is an even number of 1's

no errors:

1	0	1	0	1	1
1	1	1	1	0	0
0	1	1	1	0	1
1	0	1	0	1	0

## two-dimensional bit parity:

- detect *and correct* single bit errors



**detected and correctable single-bit error:**

1	0	1	0	1	1
<del>1</del>	<del>0</del>	<del>1</del>	<del>1</del>	<del>0</del>	<del>0</del>
0	1	1	1	0	1
1	0	1	0	1	0

$\downarrow$  parity error  
 $\rightarrow$  parity error

# UDP/TCP/IP checksum (review)

*Goal:* detect errors (*i.e.*, flipped bits) in transmitted segment

## sender:

- treat contents of UDP segment (including UDP header fields and IP addresses) as sequence of 16-bit integers
- **checksum:** addition (one's complement sum) of segment content
- checksum value put into UDP checksum field

## receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
  - not equal - error detected
  - equal - no error detected. *But maybe errors nonetheless?* More later ....

# Cyclic Redundancy Check (CRC)


- error-detection code more powerful than TCP/UDP/IP's checksum
  - used by Ethernet and WiFi
  - can detect all burst errors less than  $r+1$  bits

- both sender and receiver know  $G$  in advance:

- $G$ : generator of length  $r+1$  bits

- sender computes the CRC bits  $R$  and transmits  $\langle D, R \rangle$

- $D$ : data bits
- $R = D \cdot 2^r \% G = \langle D, 00 \dots 0 \rangle \% G$  is the remainder
  - the CRC bits  $R$  is of length  $r$  bits
  - use **bitwise-XOR** for addition/subtraction
  - $\langle D, R \rangle$  is divisible by  $G$



- receiver checks out whether  $\langle D, R \rangle \% G = 0$

- non-zero remainder  $\rightarrow$  error detected
- zero remainder  $\rightarrow$  either no error or error not detected


# Cyclic Redundancy Check (CRC): example

■  $G = 1001$  and  $D = 101110$

- in binary representations
- $r = 3$  in this example

■ sender computes the remainder  $R$

- $R = D \cdot 2^r \% G$   
 $= \langle D, 00 \dots 0 \rangle \% G$



# Link layer, LANs: roadmap

- introduction
- error detection, correction
- **multiple access protocols**
- LANs
  - addressing, ARP
  - Ethernet
  - switches
  - VLANs
- link virtualization: MPLS
- data center networking



- a day in the life of a web request

# Multiple access links, protocols

two types of “links”:

- point-to-point
  - point-to-point link between Ethernet switch and host
- broadcast (shared wire or medium)
  - old-fashioned Ethernet
  - 802.11 wireless LAN, mobile networks (4G/5G), satellite




shared wire (e.g.,  
cabled Ethernet)



shared radio: 4G/5G



shared radio: WiFi



shared radio: satellite



humans at a cocktail party  
(shared air, acoustical)



# Multiple access protocols

- single shared broadcast channel
- two or more simultaneous transmissions by nodes: interference
  - *collision* if node receives two or more signals at the same time

## multiple access protocol

- distributed (or centralized) algorithm that determines how nodes share channel, i.e., determine when node can transmit
- communication about channel sharing must use channel itself!
  - no out-of-band channel for coordination

# An ideal multiple access protocol

*given:* multiple access channel (MAC) of capacity  $R$  bps

*desiderate:*

1. when one node wants to transmit, it can send at rate  $R$ .
2. when  $M$  nodes want to transmit, each can send at average rate  $R/M$
3. fully decentralized:
  - no special node to coordinate transmissions
    - no single point of failure
  - no synchronization of clocks, slots
4. simple

# MAC protocols: taxonomy

three broad classes:

- **channel partitioning**

- divide channel into smaller “pieces” (time slots, frequency, ...)
- allocate piece to node/user for **exclusive** use

- ***random access***

- channel not divided, allow collisions
- “recover” from collisions


- **“taking turns”**

- nodes take turns, but nodes with more to send can take longer turns
  - similar to round robin

# Channel partitioning MAC protocols: TDMA

## TDMA: time division multiple access

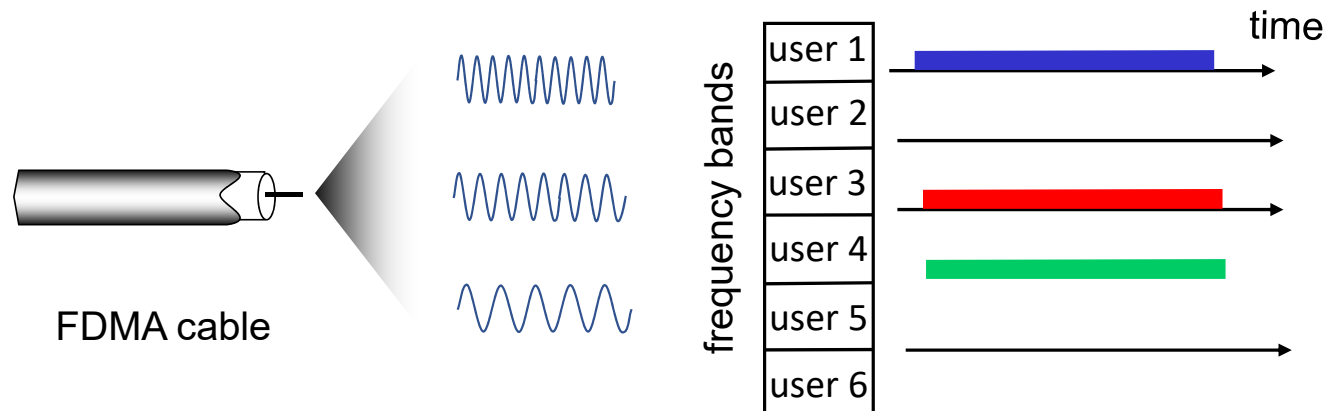
- access to channel in “rounds”
- each station gets fixed length slot (length = packet transmission time) in each round
- unused slots go idle
- example: 6-station LAN, 1,3,4 have packets to send, slots 2,5,6 idle



# Channel partitioning MAC protocols: FDMA

## FDMA: frequency division multiple access

- channel spectrum divided into frequency bands
- each station assigned fixed frequency band
- unused transmission time in frequency bands go idle
- example: 6-station LAN, 1,3,4 have packet to send, frequency bands 2,5,6 idle



# Random access protocols

- when node has packet to send
  - transmit at full channel data rate  $R$ .
  - no *a priori* coordination among nodes
- two or more transmitting nodes: “collision”
- **random access MAC protocol** specifies:
  - how to detect collisions
  - how to recover from collisions (e.g., via delayed retransmissions)
- examples of random access MAC protocols:
  - ALOHA, slotted ALOHA
  - CSMA, CSMA/CD, CSMA/CA

# Slotted ALOHA

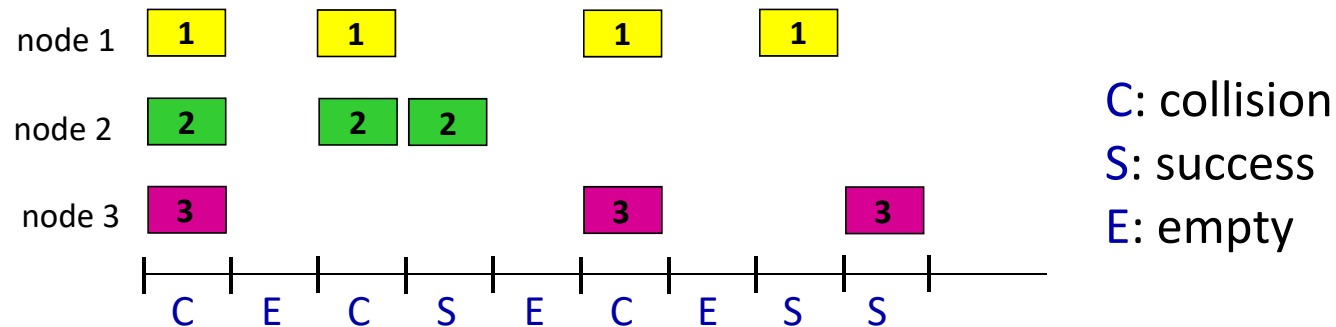
## assumptions:

- time is divided into equal-sized slots
  - slot: time to transmit 1 frame
    - all frames have the same size
  - nodes are time-synchronized
  - nodes start to transmit only slot beginning
- if 2 or more nodes transmit in a slot, all nodes detect collision

## operation:

- stop-and-wait
- when node obtains fresh frame, transmits in next slot
  - *if no collision*: node can send new frame in next slot
  - *if collision*: node retransmits frame in each subsequent slot with probability  $p$  until success

# Slotted ALOHA



## Pros:

- single active node can continuously transmit at full rate of channel
- highly decentralized: only slots in nodes need to be in sync
- simple

## Cons:

- collisions
- idle slots
- nodes may be able to detect collision in less than time to transmit packet
- clock synchronization



## Slotted ALOHA: efficiency

**efficiency:** long-run fraction of successful slots

- *suppose:* Given  $N$  nodes, each node has many frames to send and transmits in every slot with probability  $p$ 
  - probability that a given node has success in a slot
  - probability that *any* node has a success =  $Np(1 - p)^{N-1}$

## Slotted ALOHA: efficiency (cont'd)

- slotted ALOHA's efficiency = probability that *any* node has a success
  - Given  $N$  and  $p$ ,  $f(p) = Np(1 - p)^{N-1}$
- max efficiency: find  $p^*$  that maximizes  $f(p)$


- for many nodes, take limit of  $Np^*(1 - p^*)^{N-1}$  as  $N$  goes to infinity:

$$\lim_{N \rightarrow \infty} f(p^*) = \lim_{N \rightarrow \infty} N \cdot \frac{1}{N} \cdot \left(1 - \frac{1}{N}\right)^{N-1} = \lim_{N \rightarrow \infty} \left(\frac{1}{1 + \frac{1}{N-1}}\right)^{N-1} = \frac{1}{e}$$

- slotted ALOHA's efficiency is at best  $1/e = 37\%$ 
  - 37% of time is useful transmissions.

# Pure ALOHA

- unslotted Aloha: simpler, no synchronization
  - when frame first arrives: transmit immediately
- collision probability increases with no synchronization:
  - frame sent at  $t_0$  collides with other frames sent in  $[t_0-1, t_0+1]$



- pure Aloha efficiency: 18% !


# CSMA (carrier sense multiple access)

simple **CSMA**: listen before transmit:

- if channel sensed **idle**: transmit entire frame
- if channel sensed **busy**: defer transmission
- human analogy: don't interrupt others!

# CSMA: collisions

- collisions *can* still occur if two nodes send frames around the same time
  - propagation delay means two nodes may not hear each other's just-started transmission
- **collision**: entire packet transmission time wasted
  - distance & propagation delay play role in determining collision probability




# CSMA/CD

**CSMA/CD:** CSMA with *collision detection*

- collisions *detected* within short time
  - colliding transmissions aborted, reducing channel wastage
  - collision detection easy in wired, difficult with wireless
- human analogy: the polite conversationalist

# CSMA/CD

- CSMA/CS reduces the amount of time wasted in collisions
  - transmission aborted on collision detection



# Ethernet CSMA/CD algorithm

1. NIC receives datagram from network layer, creates frame
2. NIC senses channel:
  - if **idle**: start frame transmission.
  - if **busy**: wait until channel idle, then transmit
3. If NIC transmits entire frame without collision, NIC is done with frame !
4. If NIC detects another transmission while sending: abort, send jam signal
5. After aborting, NIC enters *binary (exponential) backoff*:
  - after  $m$ -th collision, NIC chooses  $K$  at random from  $\{0, 1, 2, \dots, 2^m - 1\}$ . NIC waits  $K$  slot times, returns to Step 2. (1 slot = minimum frame size = 64 bytes for 10/100 Mbps or 512 bytes for 1Gbps)
  - more collisions: longer backoff interval



## CSMA/CD efficiency

- $t_{prop} \triangleq$  maximum propagation delay between 2 nodes in LAN
- $t_{trans} \triangleq$  time to transmit max-size frame
- efficiency is (approximately)

$$\rho \doteq \frac{1}{1 + 5 \cdot \frac{t_{prop}}{t_{trans}}}$$

- efficiency goes to 1 as  $t_{prop}/t_{trans}$  goes to 0
- CSMA/CD
  - better efficiency than ALOHA
  - simple, cheap, decentralized

# “Taking turns” MAC protocols

## channel partitioning MAC protocols:

- share channel *efficiently* and *fairly* at high load
- inefficient at low load:  $1/N$  bandwidth allocated even if only 1 active node!

## random access MAC protocols

- efficient at low load: single node can fully utilize channel
- high load: collision overhead


## “taking turns” protocols

- look for best of both worlds!

# “Taking turns” MAC protocols

## polling:


- master node “invites” other nodes to transmit in turn
- typically used with “dumb” devices
- concerns:
  - polling overhead
  - latency
  - single point of failure (master)



# “Taking turns” MAC protocols

## token passing:

- control *token* passed from one node to next sequentially.
- token message
- concerns:
  - token overhead
  - latency
  - single point of failure (token)



# Summary of MAC protocols

- **channel partitioning**, by time, frequency or code
  - Time Division, Frequency Division
- **random access** (dynamic),
  - ALOHA, S-ALOHA, CSMA, CSMA/CD
  - carrier sensing: easy in some technologies (wired), hard in others (wireless)
  - CSMA/CD used in Ethernet
  - CSMA/CA used in 802.11
- **taking turns**
  - polling from central site, token passing
  - Bluetooth, FDDI, token ring

# Link layer, LANs: roadmap

- introduction
- error detection, correction
- multiple access protocols
- **LANs**
  - addressing, ARP
  - Ethernet
  - switches
  - VLANs
- link virtualization: MPLS
- data center networking




- a day in the life of a web request

# MAC addresses

- IP address:
  - *network-layer* address for interface
    - used for layer-3 (network layer) forwarding
      - switch doesn't have any IP address
  - 32-bit (in IPv4)
    - e.g.: 128.119.40.136
- MAC (or LAN or physical or Ethernet) address:
  - function: used “locally” to get frame from one interface to another physically-connected interface in link layer (within same subnet)
  - 48-bit MAC address (for most LANs) burned in NIC ROM, also sometimes software settable
    - e.g.: 1A-2F-BB-76-09-AD (or 1A:2F:BB:76:09:AD)

# Each interface at host or at router

- has unique 48-bit **MAC** address
- has a locally unique 32-bit IP address (as we've seen)






# MAC addresses

- MAC address allocation administered by IEEE
- manufacturer buys portion of MAC address space (to assure uniqueness)
- analogy:
  - MAC address: like Social Security Number
  - IP address: like postal address
- MAC flat address: portability
  - can move interface from one LAN to another LAN
  - recall IP address *not* portable: depends on IP subnet to which node is attached

# ARP: address resolution protocol

*Question:* how to determine interface's MAC address, knowing its IP address?



**ARP table:** each IP (layer-3) node (host, router) on LAN has table

- IP/MAC address mappings for some LAN nodes:  
< IP address; MAC address; TTL >
- TTL (Time To Live): time after which address mapping will be forgotten (typically 20 min)

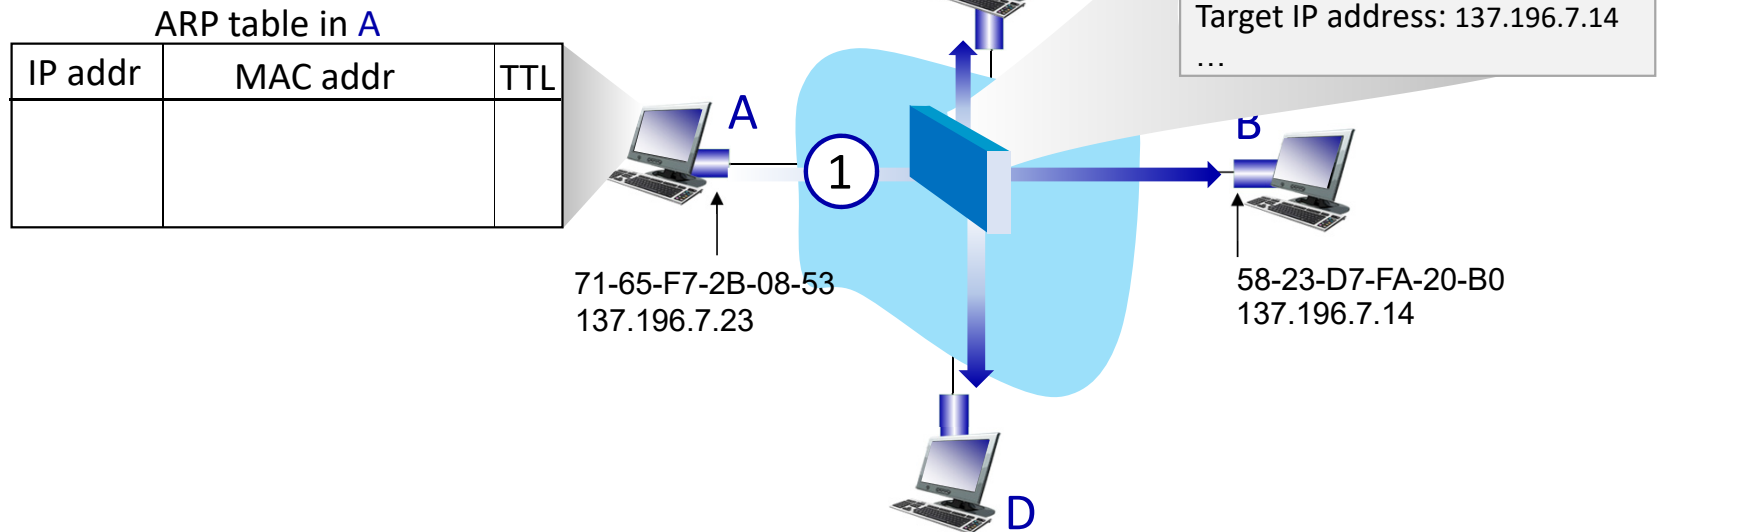
# ARP protocol in action

example: A wants to send datagram to B

- B's MAC address not in A's ARP table, so A uses ARP to find B's MAC address

A broadcasts ARP query, containing B's IP addr

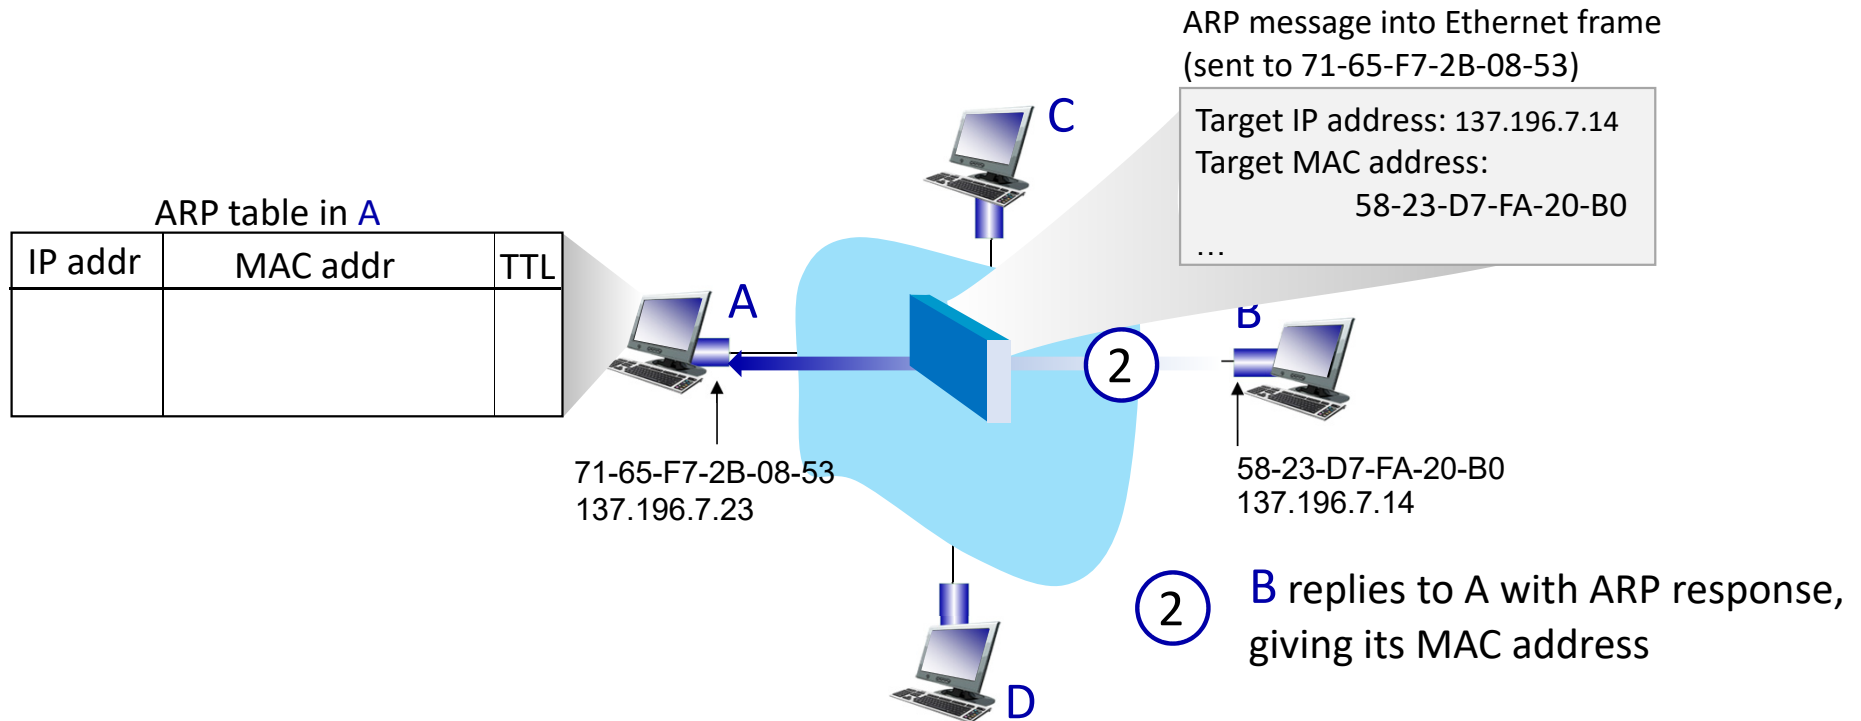
- ①
- destination MAC address = FF-FF-FF-FF-FF-FF
  - all nodes on LAN receive ARP query



# ARP protocol in action

example: A wants to send datagram to B

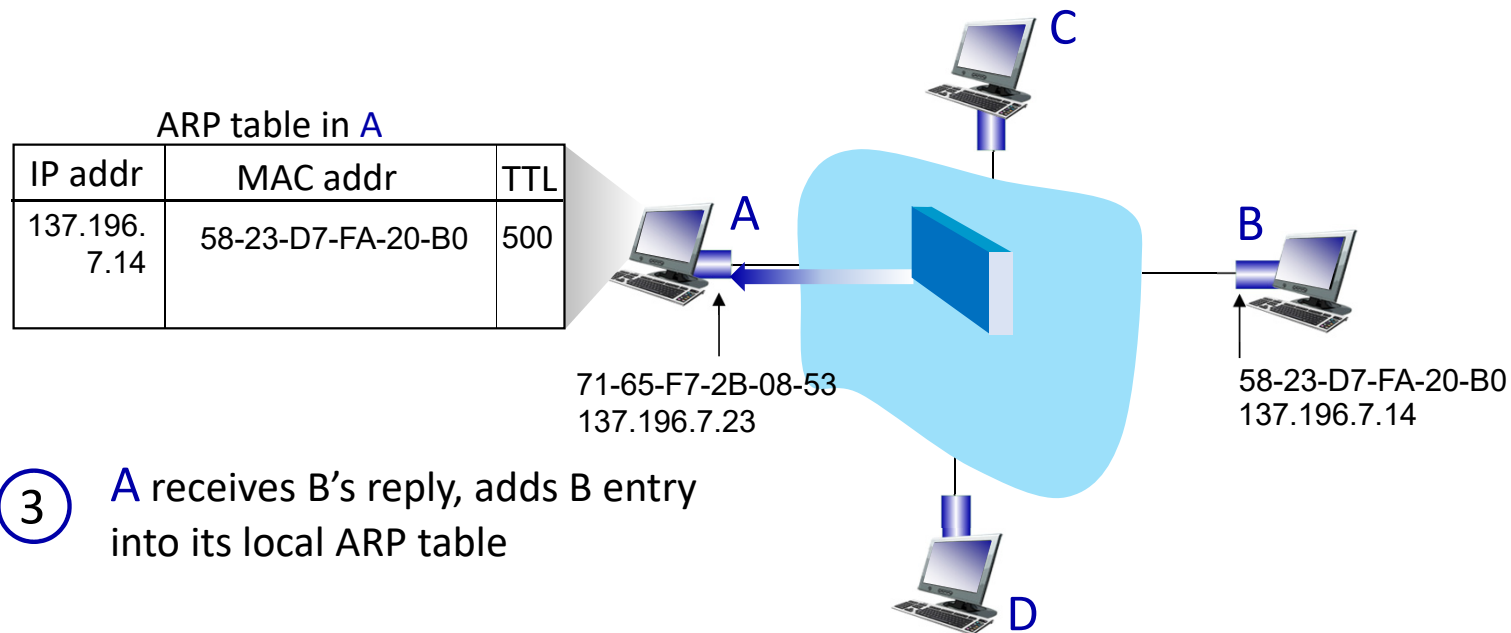
- B's MAC address not in A's ARP table, so A uses ARP to find B's MAC address



# ARP protocol in action

example: A wants to send datagram to B


- B's MAC address not in A's ARP table, so A uses ARP to find B's MAC address



# Routing to another subnet: addressing


walkthrough: sending a datagram from *A* to *B* via *R*

- focus on addressing – at IP (datagram) and MAC layer (frame) levels
- assume that:
  - A knows B's IP address
  - A knows IP address of first hop router, R (how?)
  - A knows R's MAC address (how?)




# Routing to another subnet: addressing

- A creates IP datagram with IP source A, destination B
- A creates link-layer frame containing A-to-B IP datagram
  - R's MAC address is frame's destination



# Routing to another subnet: addressing


- frame sent from A to R
- frame received at R, datagram removed, passed up to IP






# Routing to another subnet: addressing

- R determines outgoing interface, passes datagram with IP source A, destination B to link layer
- R creates link-layer frame containing A-to-B IP datagram. Frame destination address: B's MAC address




# Routing to another subnet: addressing

- R determines outgoing interface, passes datagram with IP source A, destination B to link layer
- R creates link-layer frame containing A-to-B IP datagram. Frame destination address: B's MAC address
- transmits link-layer frame



# Routing to another subnet: addressing

- B receives frame, extracts IP datagram destination B
- B passes datagram up protocol stack to IP



# Link layer, LANs: roadmap

- introduction
- error detection, correction
- multiple access protocols
- **LANs**
  - addressing, ARP
  - **Ethernet**
  - switches
  - VLANs
- link virtualization: MPLS
- data center networking




- a day in the life of a web request

# Ethernet

“dominant” wired LAN technology:



- first widely used LAN technology
- simpler, cheap
- kept up with speed race: 10 Mbps – 400 Gbps
- single chip, multiple speeds (e.g., Broadcom BCM5761)



*Metcalfe's Ethernet sketch*


# Ethernet: physical topology

- **bus:** popular through mid 90s
  - all nodes in same collision domain (or called “segment” later)
- **switched:** prevails today
  - *switch* in center (hosts are connected to switches)
  - each segment runs a (separate) Ethernet protocol
    - store-and-forward (frames are stored in a switch and then forwarded)
    - (different) segments do not collide with each other



# Ethernet frame structure

sending interface encapsulates IP datagram (or other network layer protocol packet) in **Ethernet frame**



## *preamble:*

- used to synchronize receiver, sender clock rates
- 7 bytes of 10101010 followed by one byte of 10101011

# Ethernet frame structure (more)



- **addresses:** 6 byte source, destination MAC addresses
  - if adapter receives frame with matching destination address, or with broadcast address (e.g., ARP packet), it passes data in frame to network layer protocol
  - otherwise, adapter discards frame
- **type:** indicates higher layer protocol
  - mostly IP, but others possible (e.g., Novell IPX, AppleTalk)
  - used to demultiplex up at receiver
- **CRC:** cyclic redundancy check at receiver
  - error detected: frame is dropped



# Ethernet: unreliable, connectionless

- **connectionless**: no handshaking between sending and receiving NICs
- Ethernet's MAC protocol: unslotted **CSMA/CD with binary backoff**
  - backoff and retransmit
- **unreliable**: receiving NIC doesn't send ACKs or NAKs to sending NIC
  - data in dropped frames recovered only if initial sender uses higher layer rdt (e.g., TCP), otherwise dropped data lost

# Link layer, LANs: roadmap


- introduction
- error detection, correction
- multiple access protocols
- **LANs**
  - addressing, ARP
  - Ethernet
  - **switches**
  - VLANs
- link virtualization: MPLS
- data center networking



- a day in the life of a web request

# Ethernet switch

- switch is a **link-layer** device
  - store and forward Ethernet frames
    - based on incoming frame's destination MAC address
  - when frame is to be forwarded on a segment, uses CSMA/CD to access the segment
- switch is transparent:
  - in layer 3, hosts unaware of the presence of switches
- switch is plug-and-play, self-learning
  - switches do not need to be configured




# Switch: multiple simultaneous transmissions


- hosts often have dedicated, direct connection to switch
- switches buffer packets
- Ethernet's MAC protocol used on *each* incoming link, so:
  - each link is its own collision domain
  - de facto collision-free and full-duplex
    - since 10Base-T (10Mbps rate)

# Switch: multiple simultaneous transmissions

- A-to-A' and B-to-B' can transmit simultaneously, without collisions



- but A-to-A' and C to A' *can't* happen simultaneously



# Switch forwarding table


Q: how does switch know A' reachable via interface 4, B' reachable via interface 5?

A: each switch has a **switch table**, each entry:

- (MAC address of host, interface to reach host, time stamp)
- looks like a routing/forwarding table at router!


Q: how are entries created, maintained in switch table?

- something like a routing protocol? No!



# Switch: self-learning (backward learning)

- switch *learns* which hosts can be reached through which interfaces
  - when frame received, switch “learns” location of sender: incoming LAN segment
  - records sender/location pair in switch table




MAC addr	interface	TTL
A	1	60

*Switch table  
(initially empty)*

# Self-learning, forwarding: example

- if switch table has no entry for the frame destination (e.g. A') → **flood (except the incoming link)**
- if switch table has an entry for the frame destination → **send on just one link**




MAC addr	interface	TTL
A	1	60
A'	4	60

*switch table  
(initially empty)*



# Interconnecting switches

self-learning switches can be connected together:



- Q: sending from A to G - how does S<sub>1</sub> know to forward frame destined to G via S<sub>4</sub> and S<sub>3</sub>?
- A: self learning! (works exactly the same as in single-switch case!)
    - no loop (because in a LAN, all links that cause loops are disabled)